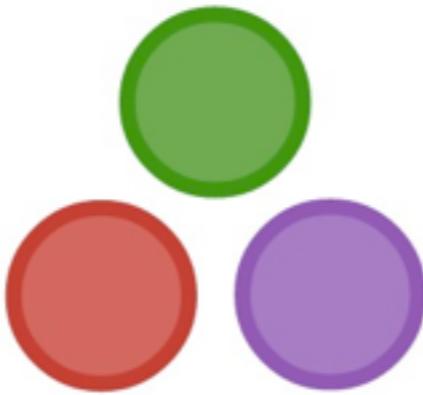


Implementing Machine Learning Algorithms in Julia

James Porter
Chicago Julia Meetup



“Julia in Anger”

James Porter
Chicago Julia Meetup

who?







OPEN SCIENCE DATA CLOUD

this talk



THE GOOD BAD AND THE UGLY

k-means clustering

perceptron

belief propagation

HMMs

eigendecomposition/SVDs

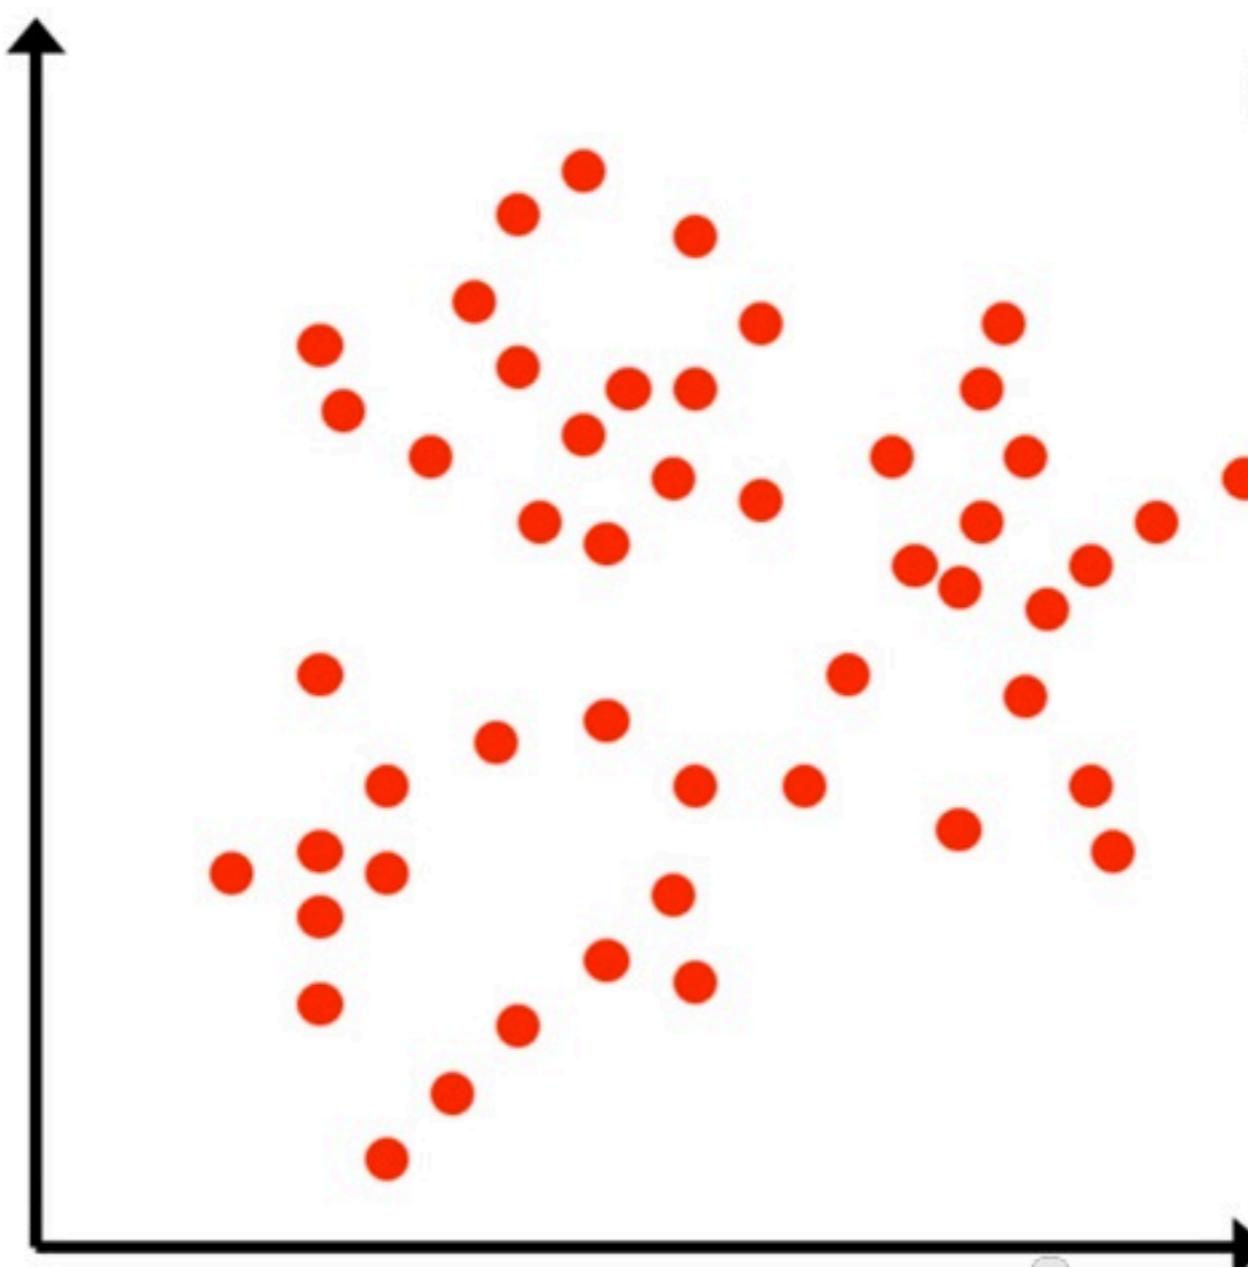
k-means clustering

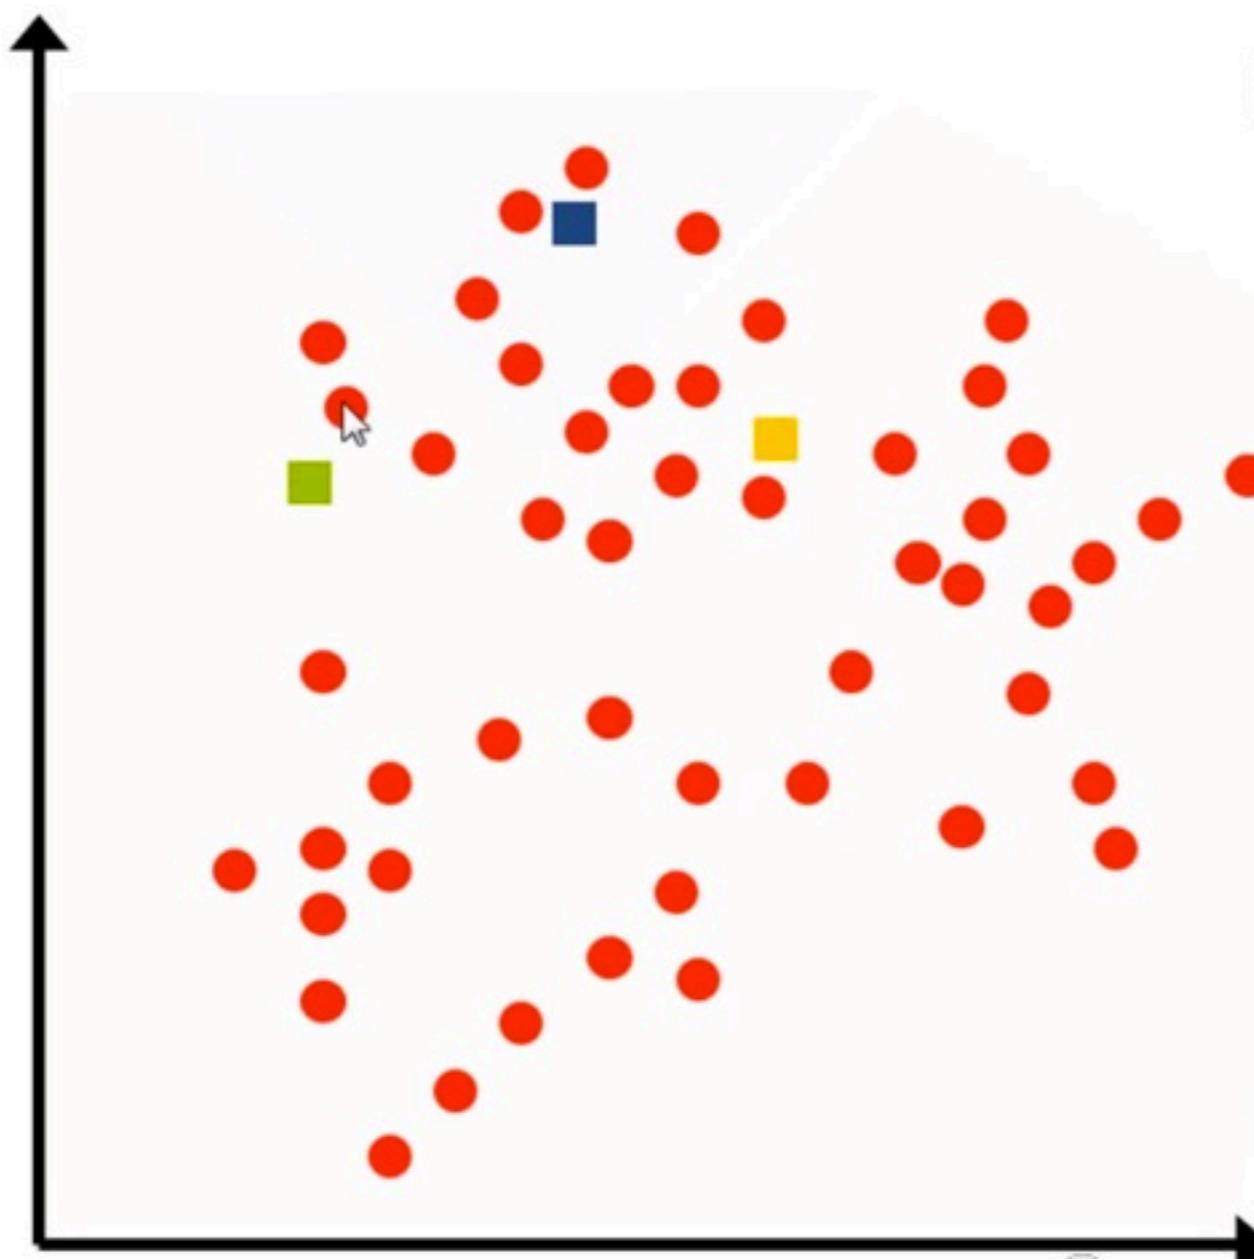
perceptron

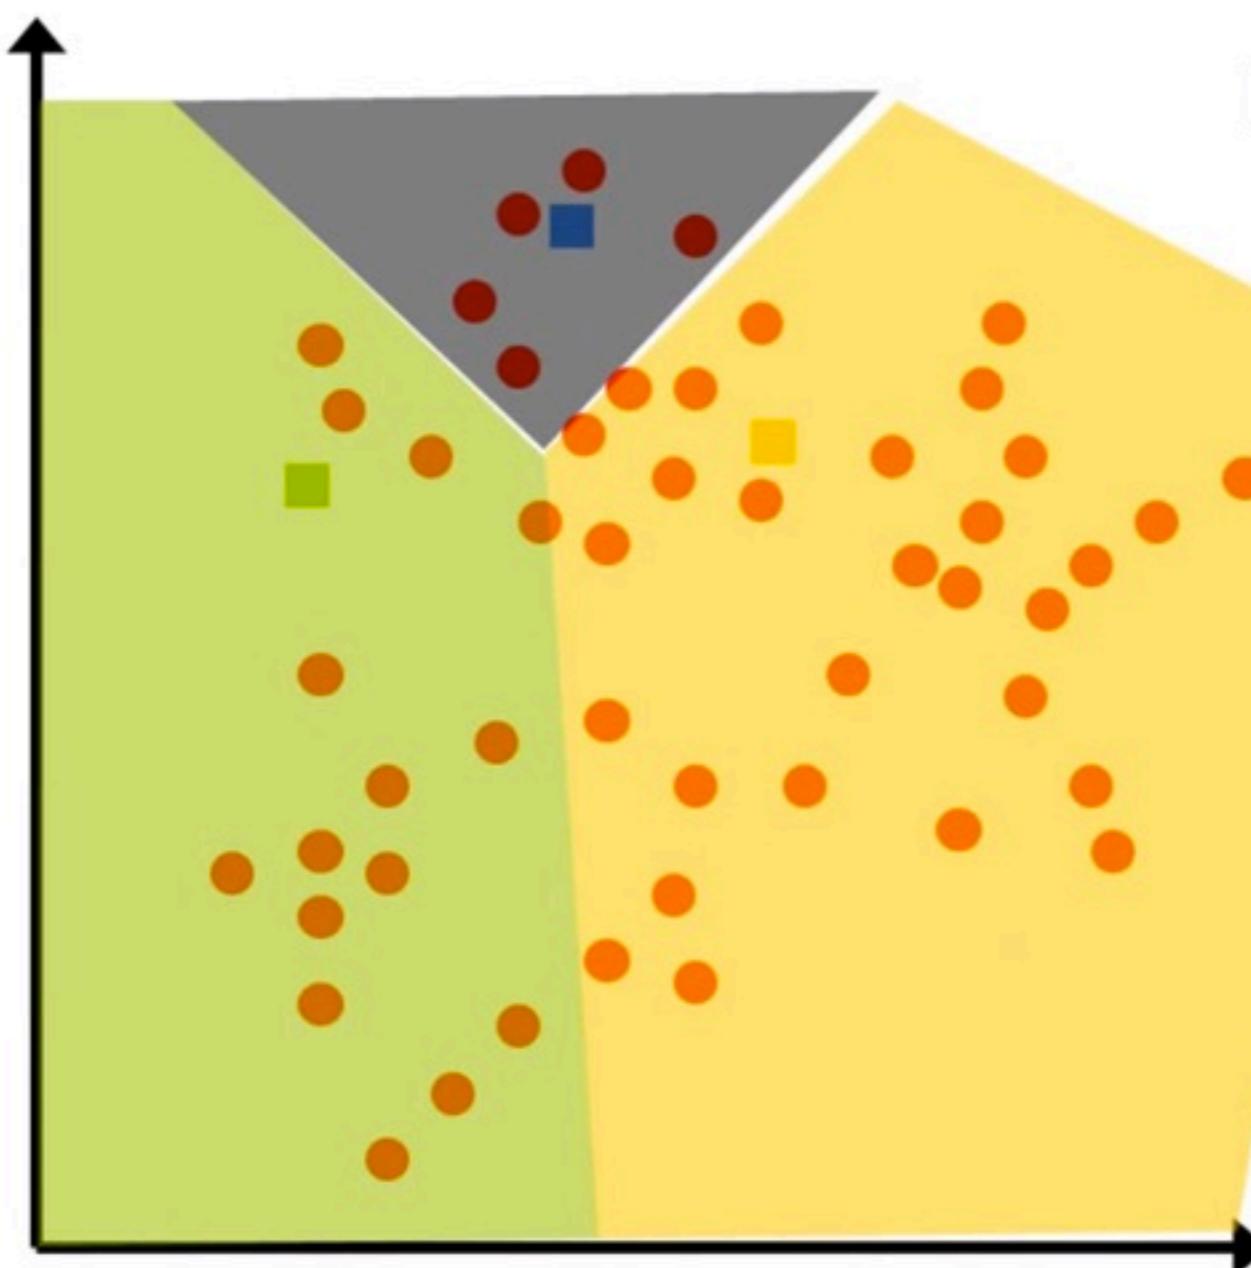
belief propagation

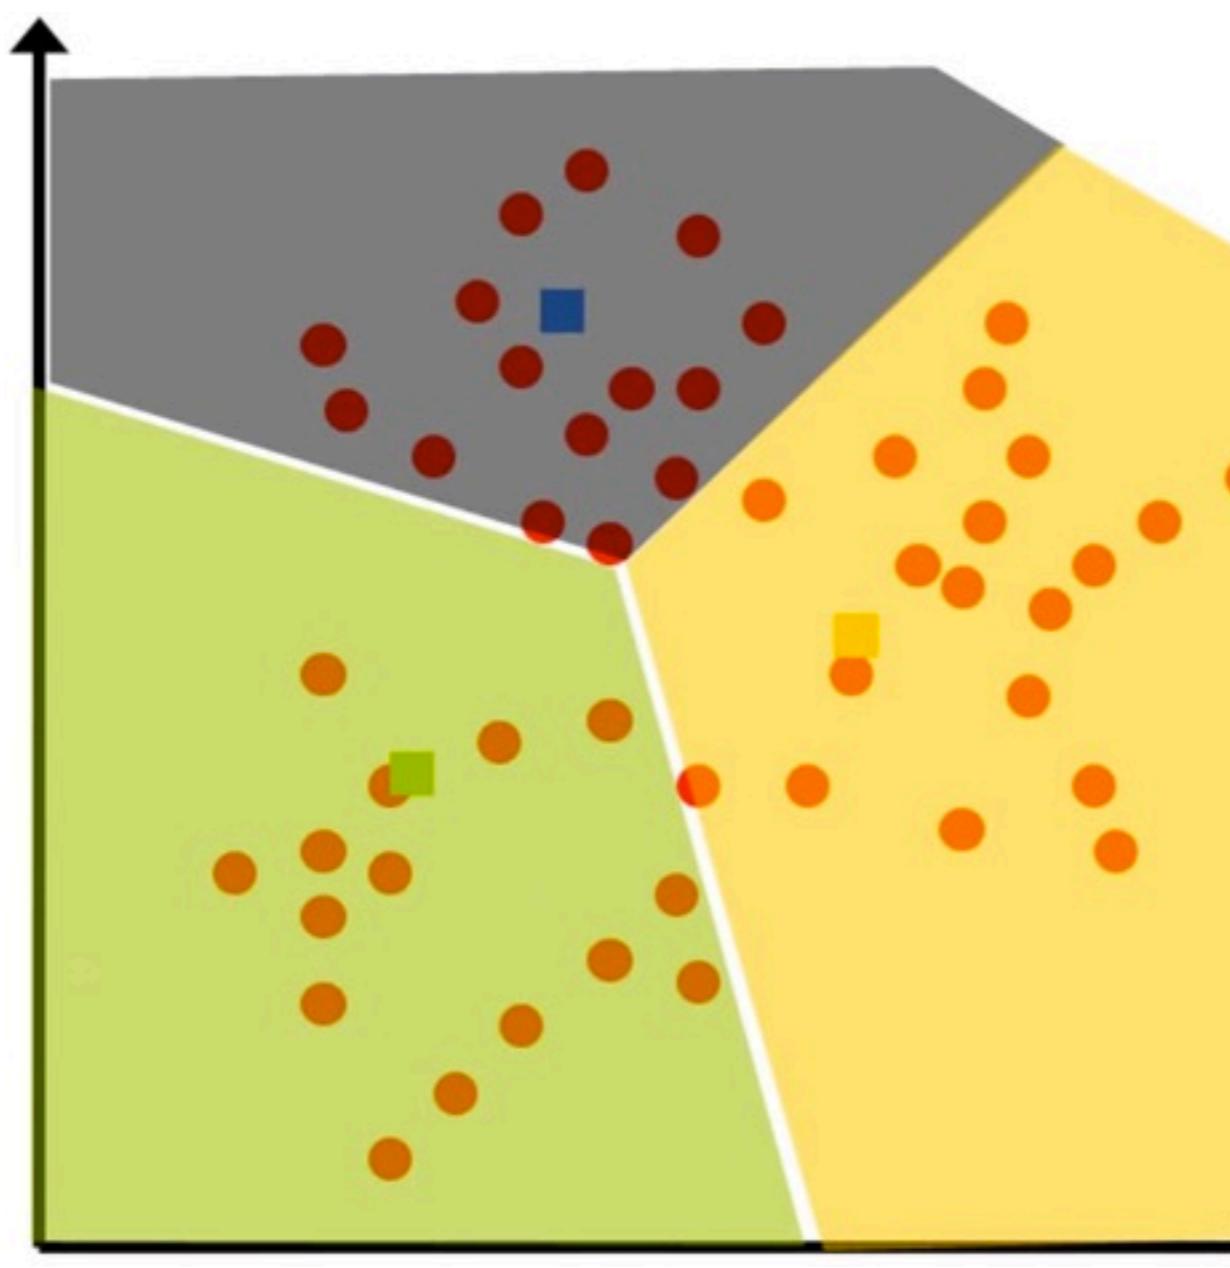
HMMs

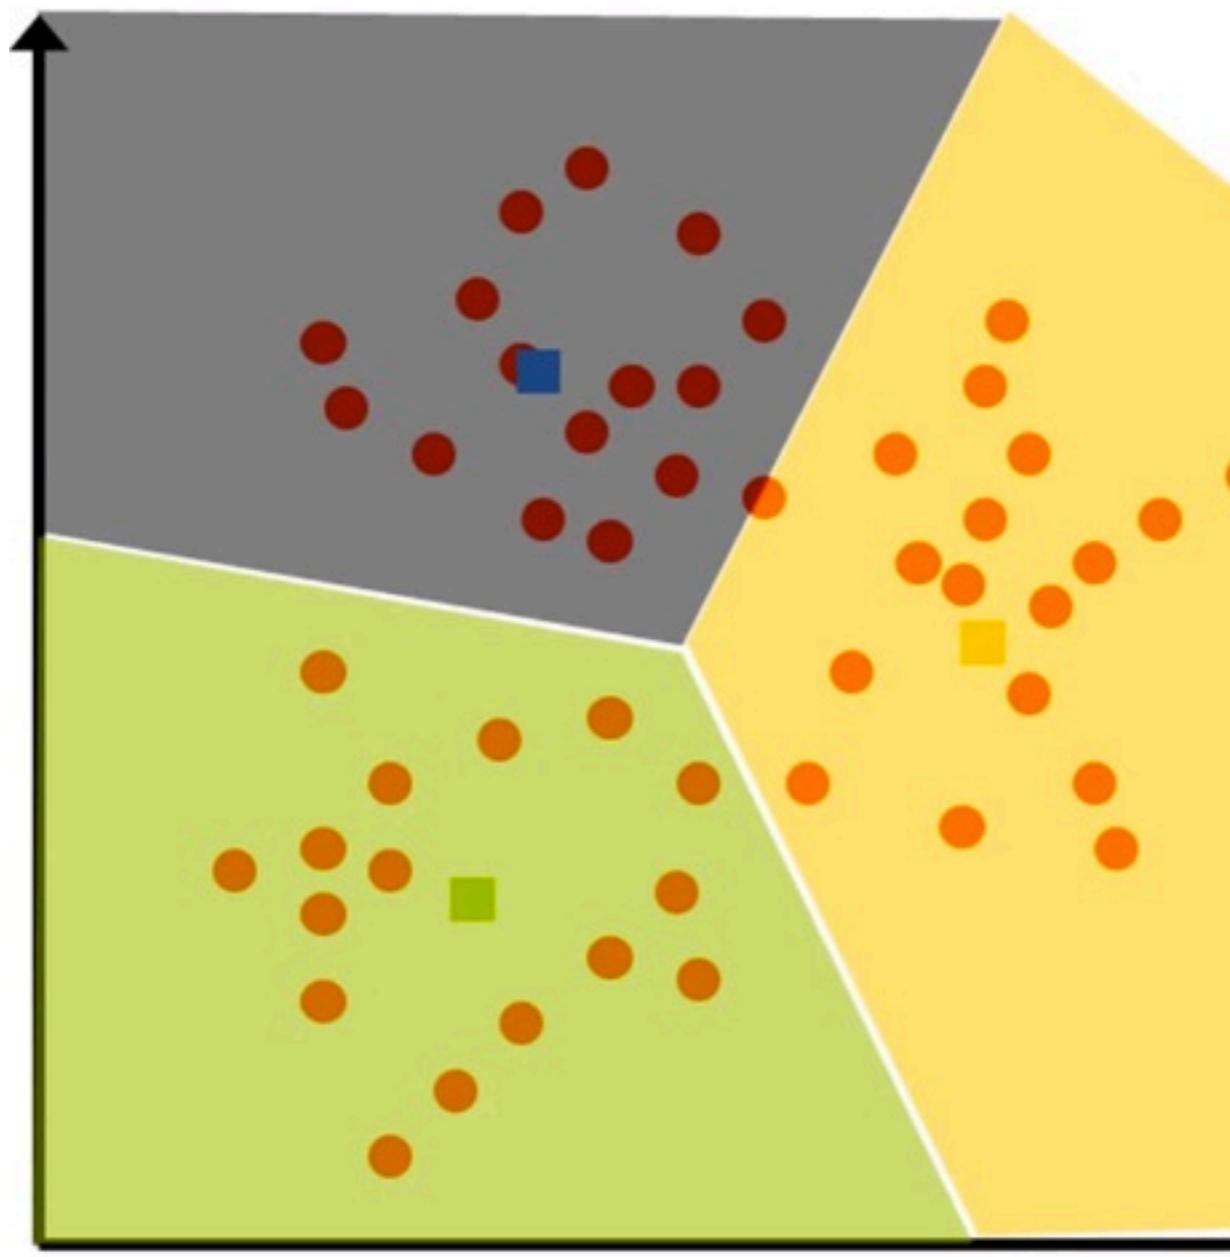
eigendecomposition/SVDs











```
# euclidean distance between vectors
function distance(x,y)
    @assert length(x) == length(y)
    res = 0
    for i = 1:length(x)
        res += (x[i]-y[i])^2
    end
    return sqrt(res)
end
```

```
# euclidean distance between vectors
function distance(x,y)
    @assert length(x) == length(y)
    res = 0
    for i = 1:length(x)
        res += (x[i]-y[i])^2
    end
    return sqrt(res)
end
```

```
julia> @elapsed distance(a,b)
0.073922868
```

```
# euclidean distance between vectors
function distance(x,y)
    @assert length(x) == length(y)
    res = 0
    for i = 1:length(x)
        res += (x[i]-y[i])^2
    end
    return sqrt(res)
end
```

```
julia> @elapsed distance(a,b)
0.073922868
```

```
# euclidean distance between vectors
function distance(x,y)
    @assert length(x) == length(y)
    res = zero(eltype(x))
    for i = 1:length(x)
        res += (x[i]-y[i])^2
    end
    return sqrt(res)
end
```

```
julia> @elapsed distance(a,b)
0.001616521
```

```
# euclidean distance between vectors
function distance(x,y)
    @assert length(x) == length(y)
    res = zero(eltype(x))
    for i = 1:length(x)
        res += (x[i]-y[i])^2
    end
    return sqrt(res)
end
```

```
julia> @elapsed distance(a,b)
0.001616521
```

```
# euclidean distance between vectors
function distance(x,y)
    @assert length(x) == length(y)
    res = zero(eltype(x))
    for i = 1:length(x)
        res += (x[i]-y[i])^2
    end
    return sqrt(res)
end
```

```
julia> @elapsed distance(a,b)
0.001616521
```

```
# euclidean distance between vectors
function distance(x,y)
    @assert length(x) == length(y)
    res = zero(eltype(x))
    @inbounds for i = 1:length(x)
        res += (x[i]-y[i])^2
    end
    return sqrt(res)
end
```

```
julia> @elapsed distance(a,b)
0.000273437
```

276x speedup!

```

while !converged && iterations < max_itr
    # Loop over all pixels and assign each to a cluster
    for i = 1:pixels
        assignments[i] = assign(data[:,i],centers)
    end

    differences = Array(Float64,size(centers,2))
    # now recalculate the centers
    for i = 1:size(centers,2)
        newcenter = center(assignments,data,i)
        differences[i] = distance(centers[:,i],newcenter)
        centers[:,i] = newcenter
    end

    if !bool(findfirst(x -> x>tol, differences))
        converged = true
    end

    iterations += 1
end

```

```
while !converged && iterations < max_itr
    # Loop over all pixels and assign each to a cluster
    for i = 1:pixels
        assignments[i] = assign(data[:,i],centers)
    end

    differences = Array(Float64,size(centers,2))
    # now recalculate the centers
    for i = 1:size(centers,2)
        newcenter = center(assignments,data,i)
        differences[i] = distance(centers[:,i],newcenter)
        centers[:,i] = newcenter
    end

    if !bool(findfirst(x -> x>tol, differences))
        converged = true
    end

    iterations += 1
end
```

```
while !converged && iterations < max_itr
    # Loop over all pixels and assign each to a cluster
    for i = 1:pixels
        assignments[i] = assign(data[:,i],centers)
    end

    differences = Array(Float64,size(centers,2))
    # now recalculate the centers
    for i = 1:size(centers,2)
        newcenter = center(assignments,data,i)
        differences[i] = distance(centers[:,i],newcenter)
        centers[:,i] = newcenter
    end

    if !bool(findfirst(x -> x>tol, differences))
        converged = true
    end

    iterations += 1
end
```

```

while !converged && iterations < max_itr
    # Loop over all pixels and assign each to a cluster
    for i = 1:pixels
        assignments[i] = assign(data[:,i],centers)
    end

    differences = Array(Float64,size(centers,2))
    # now recalculate the centers
    for i = 1:size(centers,2)
        newcenter = center(assignments,data,i)
        differences[i] = distance(centers[:,i],newcenter)
        centers[:,i] = newcenter
    end

    if !bool(findfirst(x -> x>tol, differences))
        converged = true
    end

    iterations += 1
end

```

problems?

```
for i = 1:size(centers,2)
    newcenter = center(assignments,data,i)
    centers[:,i] = newcenter
end
```

```
for i = 1:size(centers,2)
    newcenter = center(assignments,data,i)
    centers[:,i] = newcenter
end
```

```
for i = 1:size(centers,2)
    newcenter = center(assignments,data,i)
    centers[:,i] = newcenter
end
```



memory allocation!

```
for i = 1:size(centers,2)
    newcenter = center(assignments,data,i)
    centers[:,i] = newcenter
end
```

```
nextcenter = zeros(length(data[:,1]))
for i = 1:size(centers,2)
    center!(nextcenter, assignments, data, i)
    centers[:,i] = nextcenter
end
```

```
julia> @elapsed compress(img, 20)  
23.605304138
```

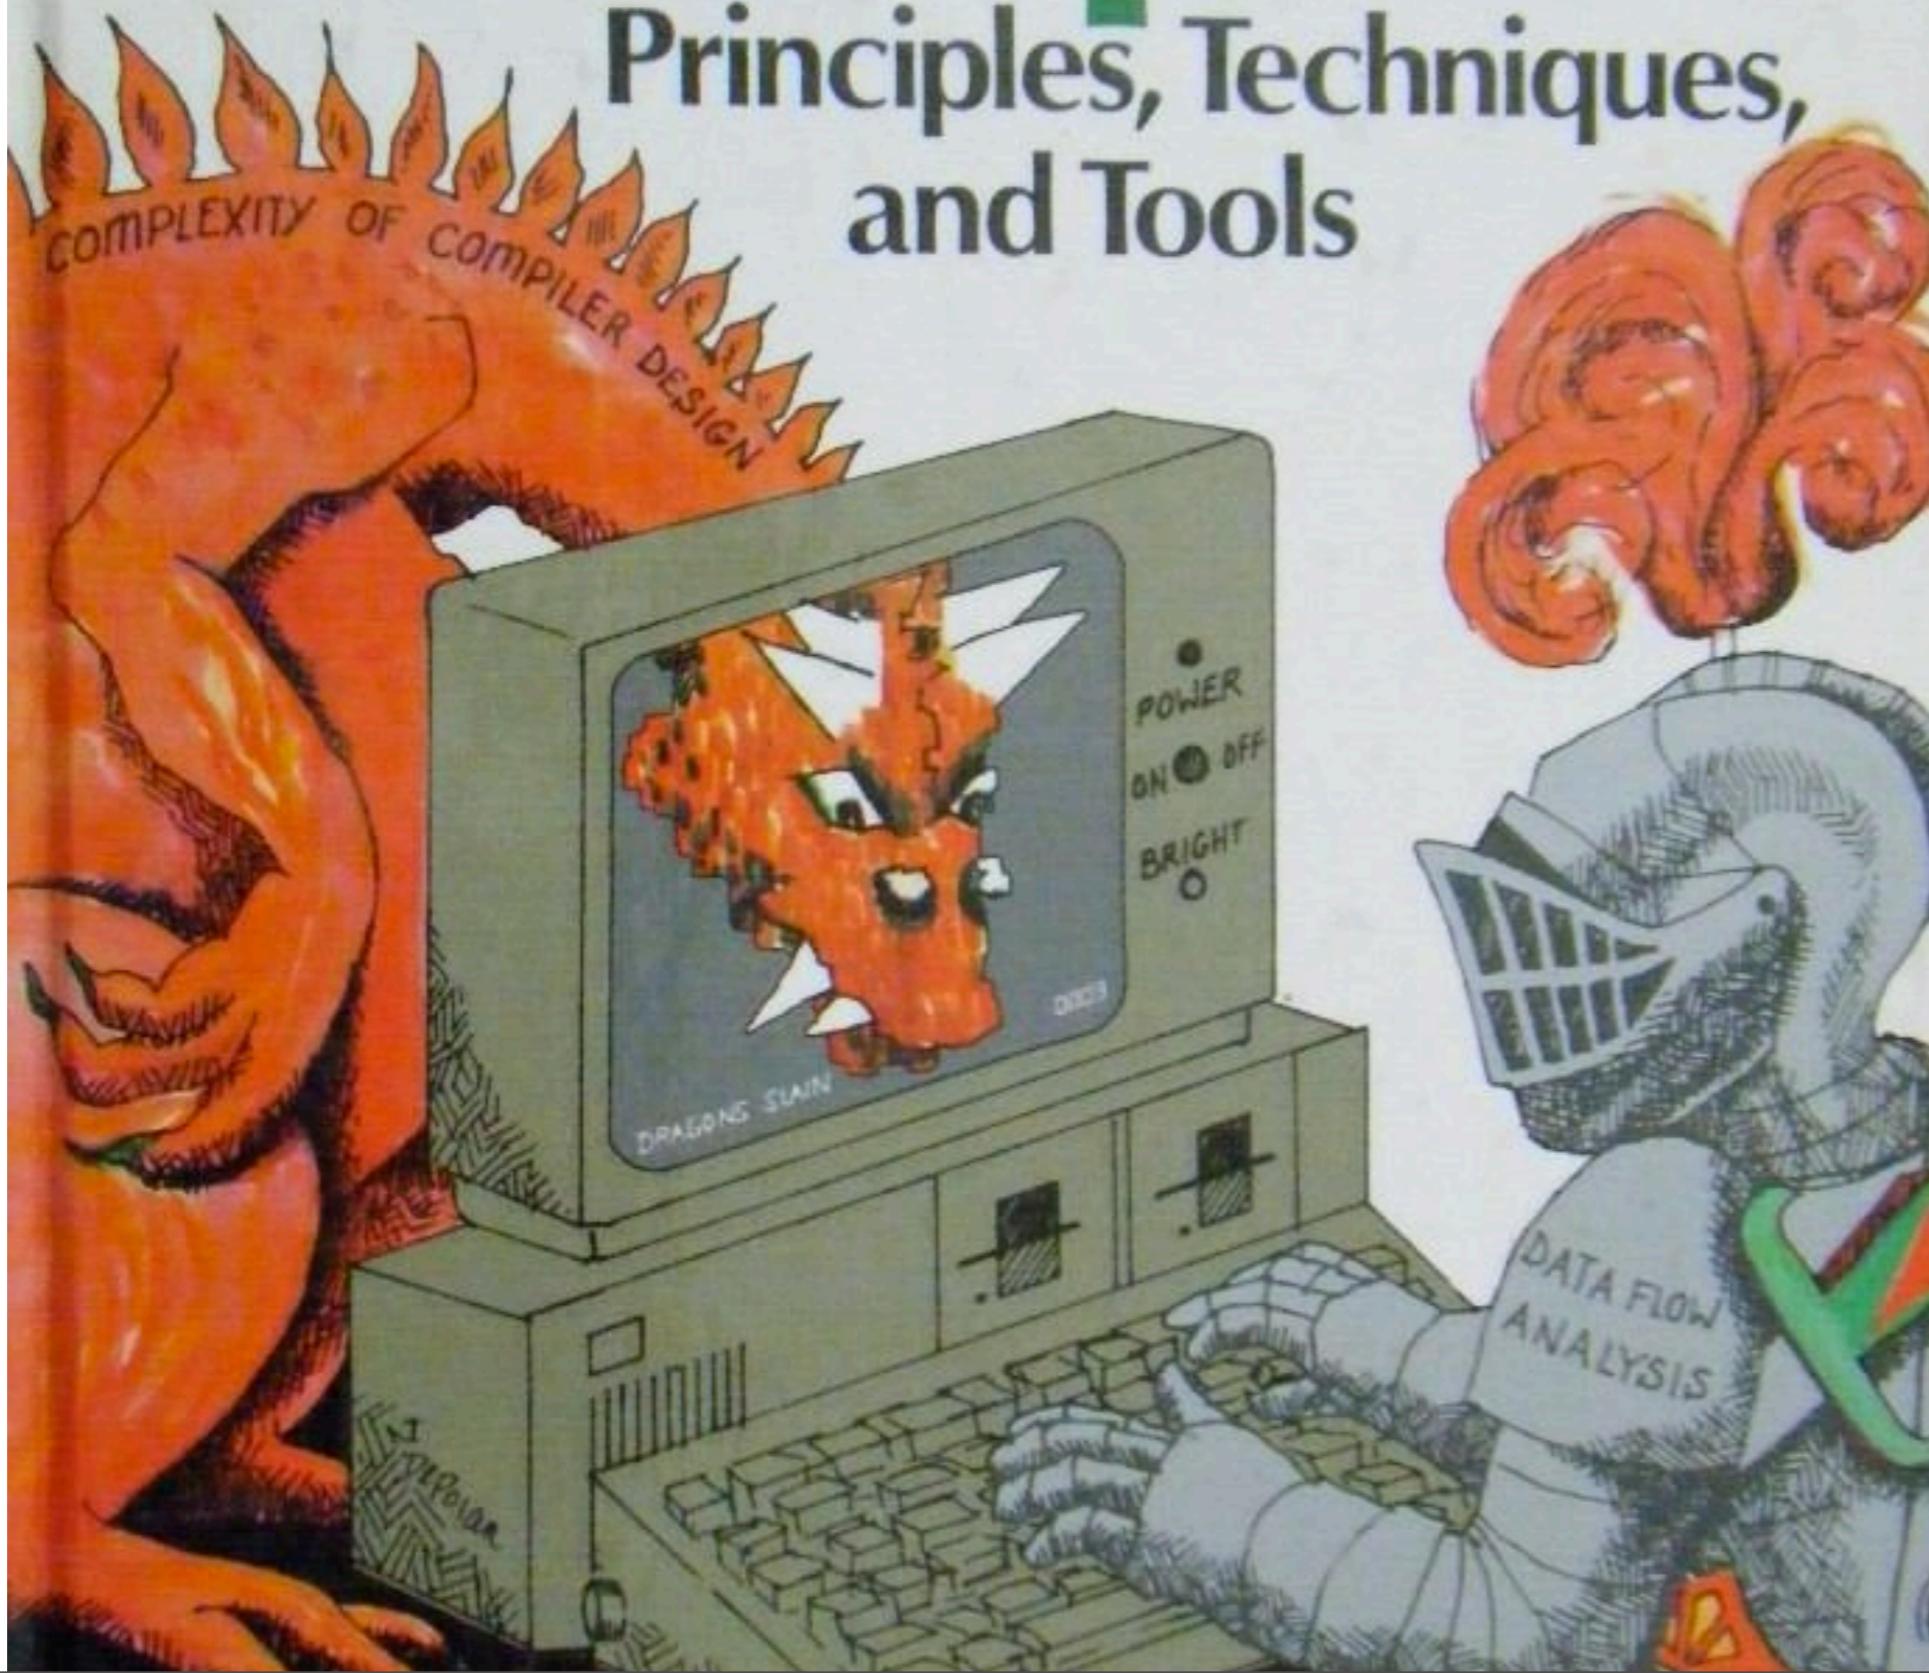
```
julia> @elapsed compress(img, 20)  
15.409054303
```

```
julia> @elapsed compress(img, 20)  
15.409054303
```

“why can’t the compiler do this for me?”

Compilers

Principles, Techniques,
and Tools



Compilers

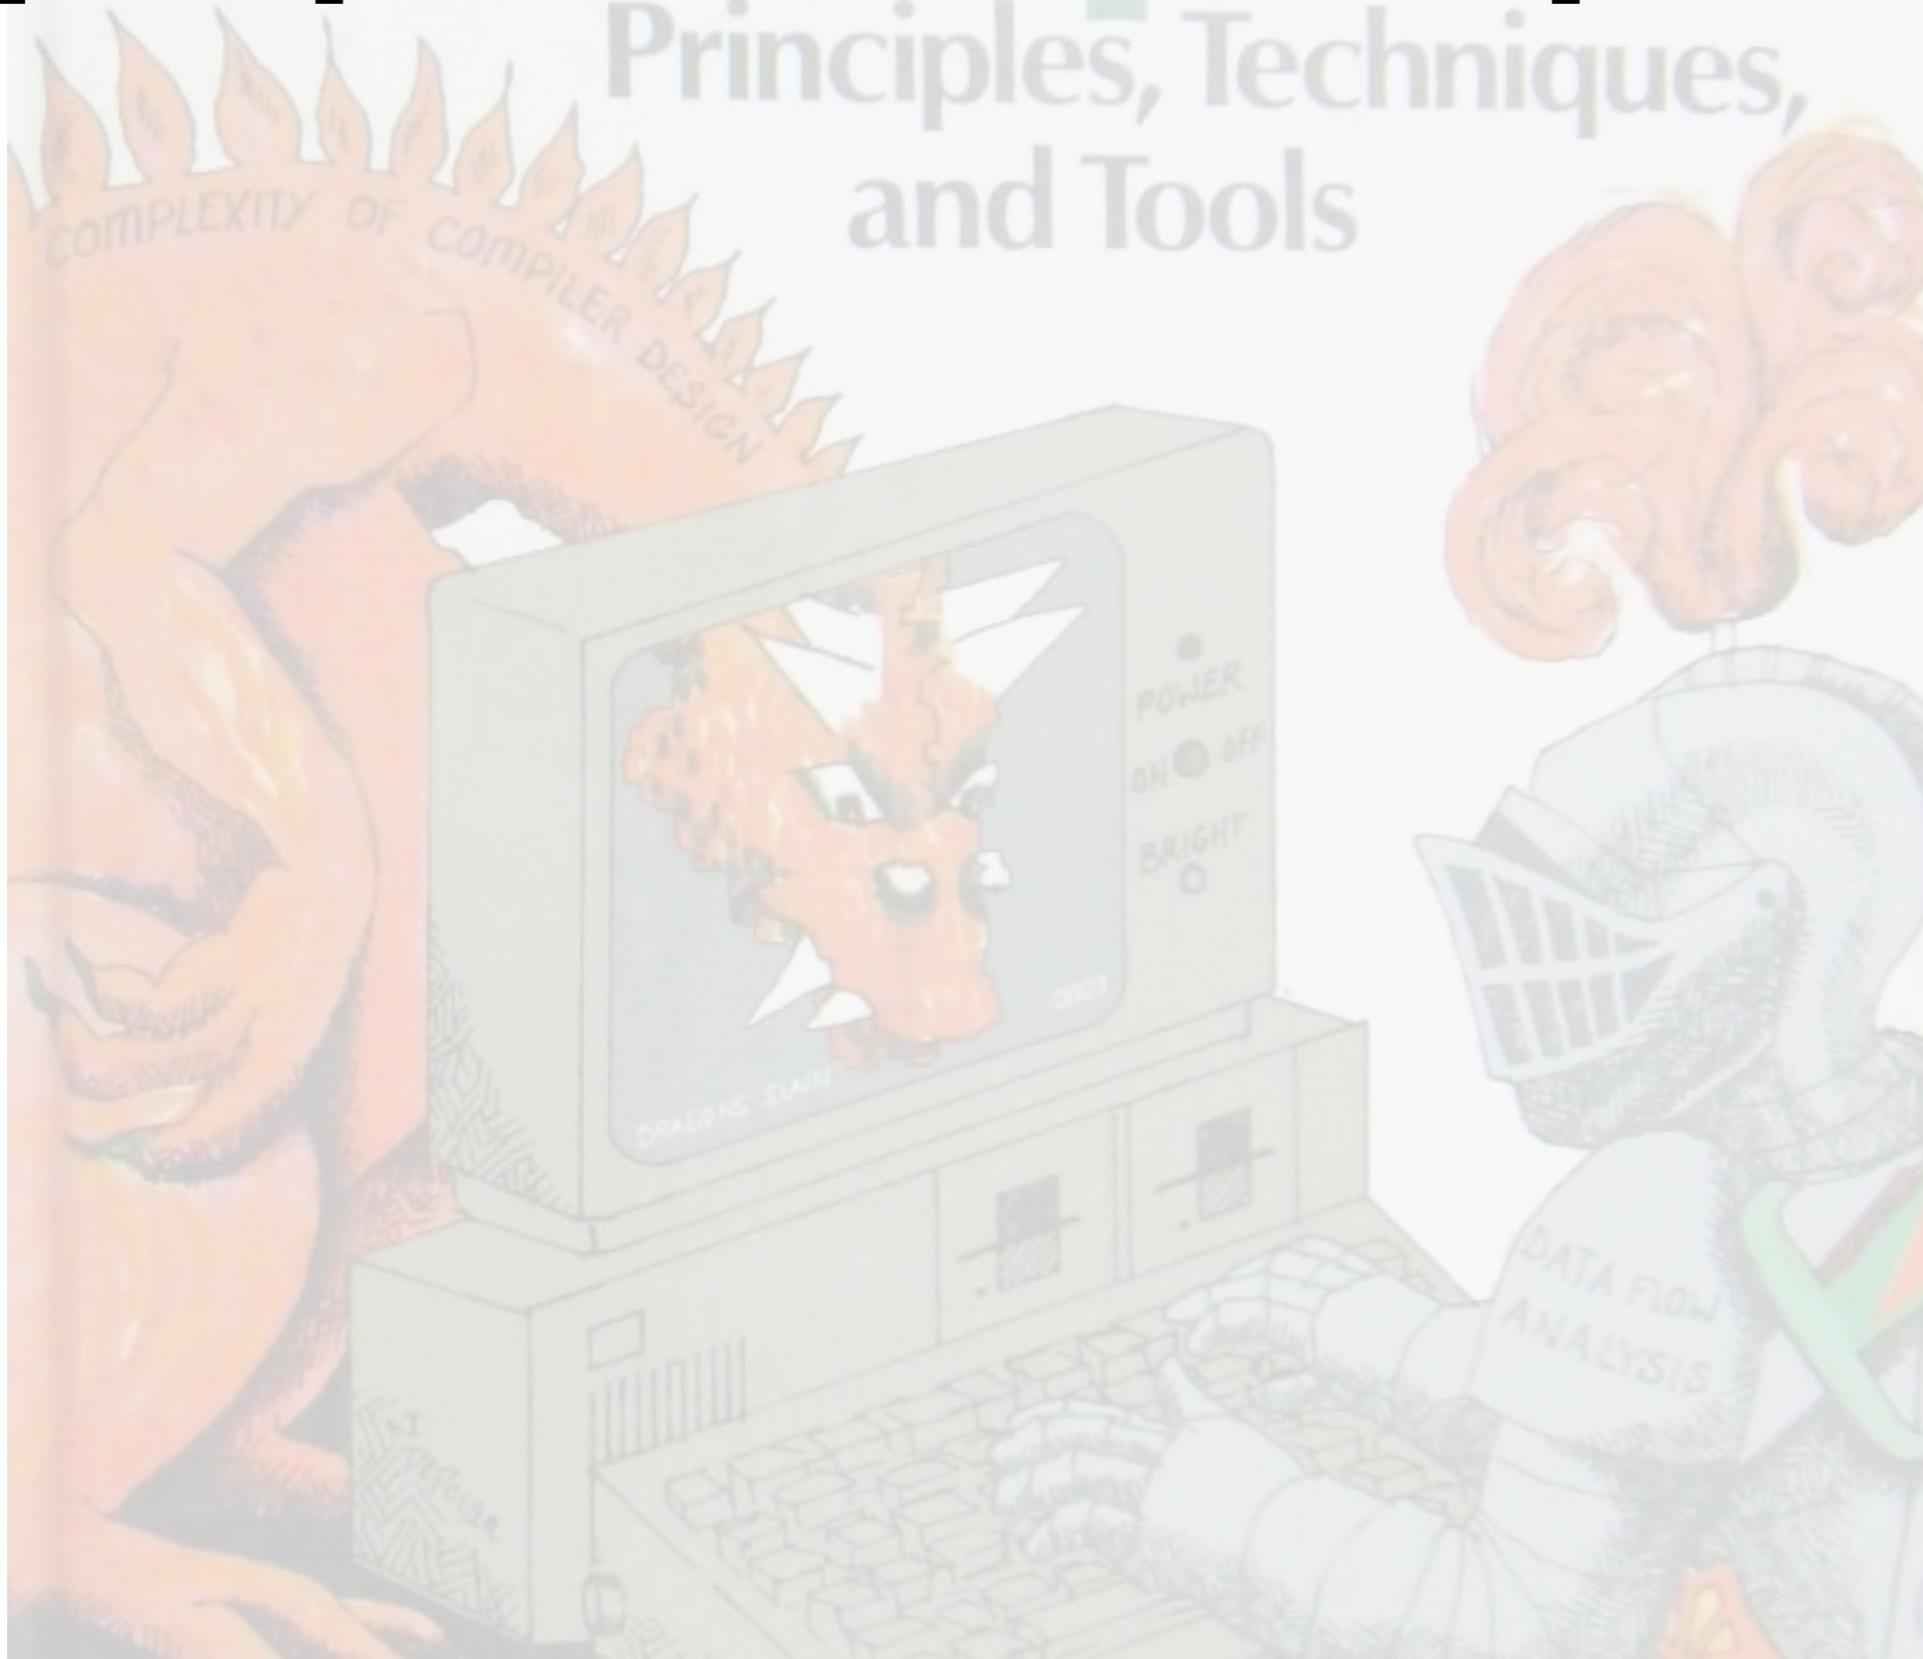
Principles, Techniques,
and Tools

two kinds



opaque

transparent

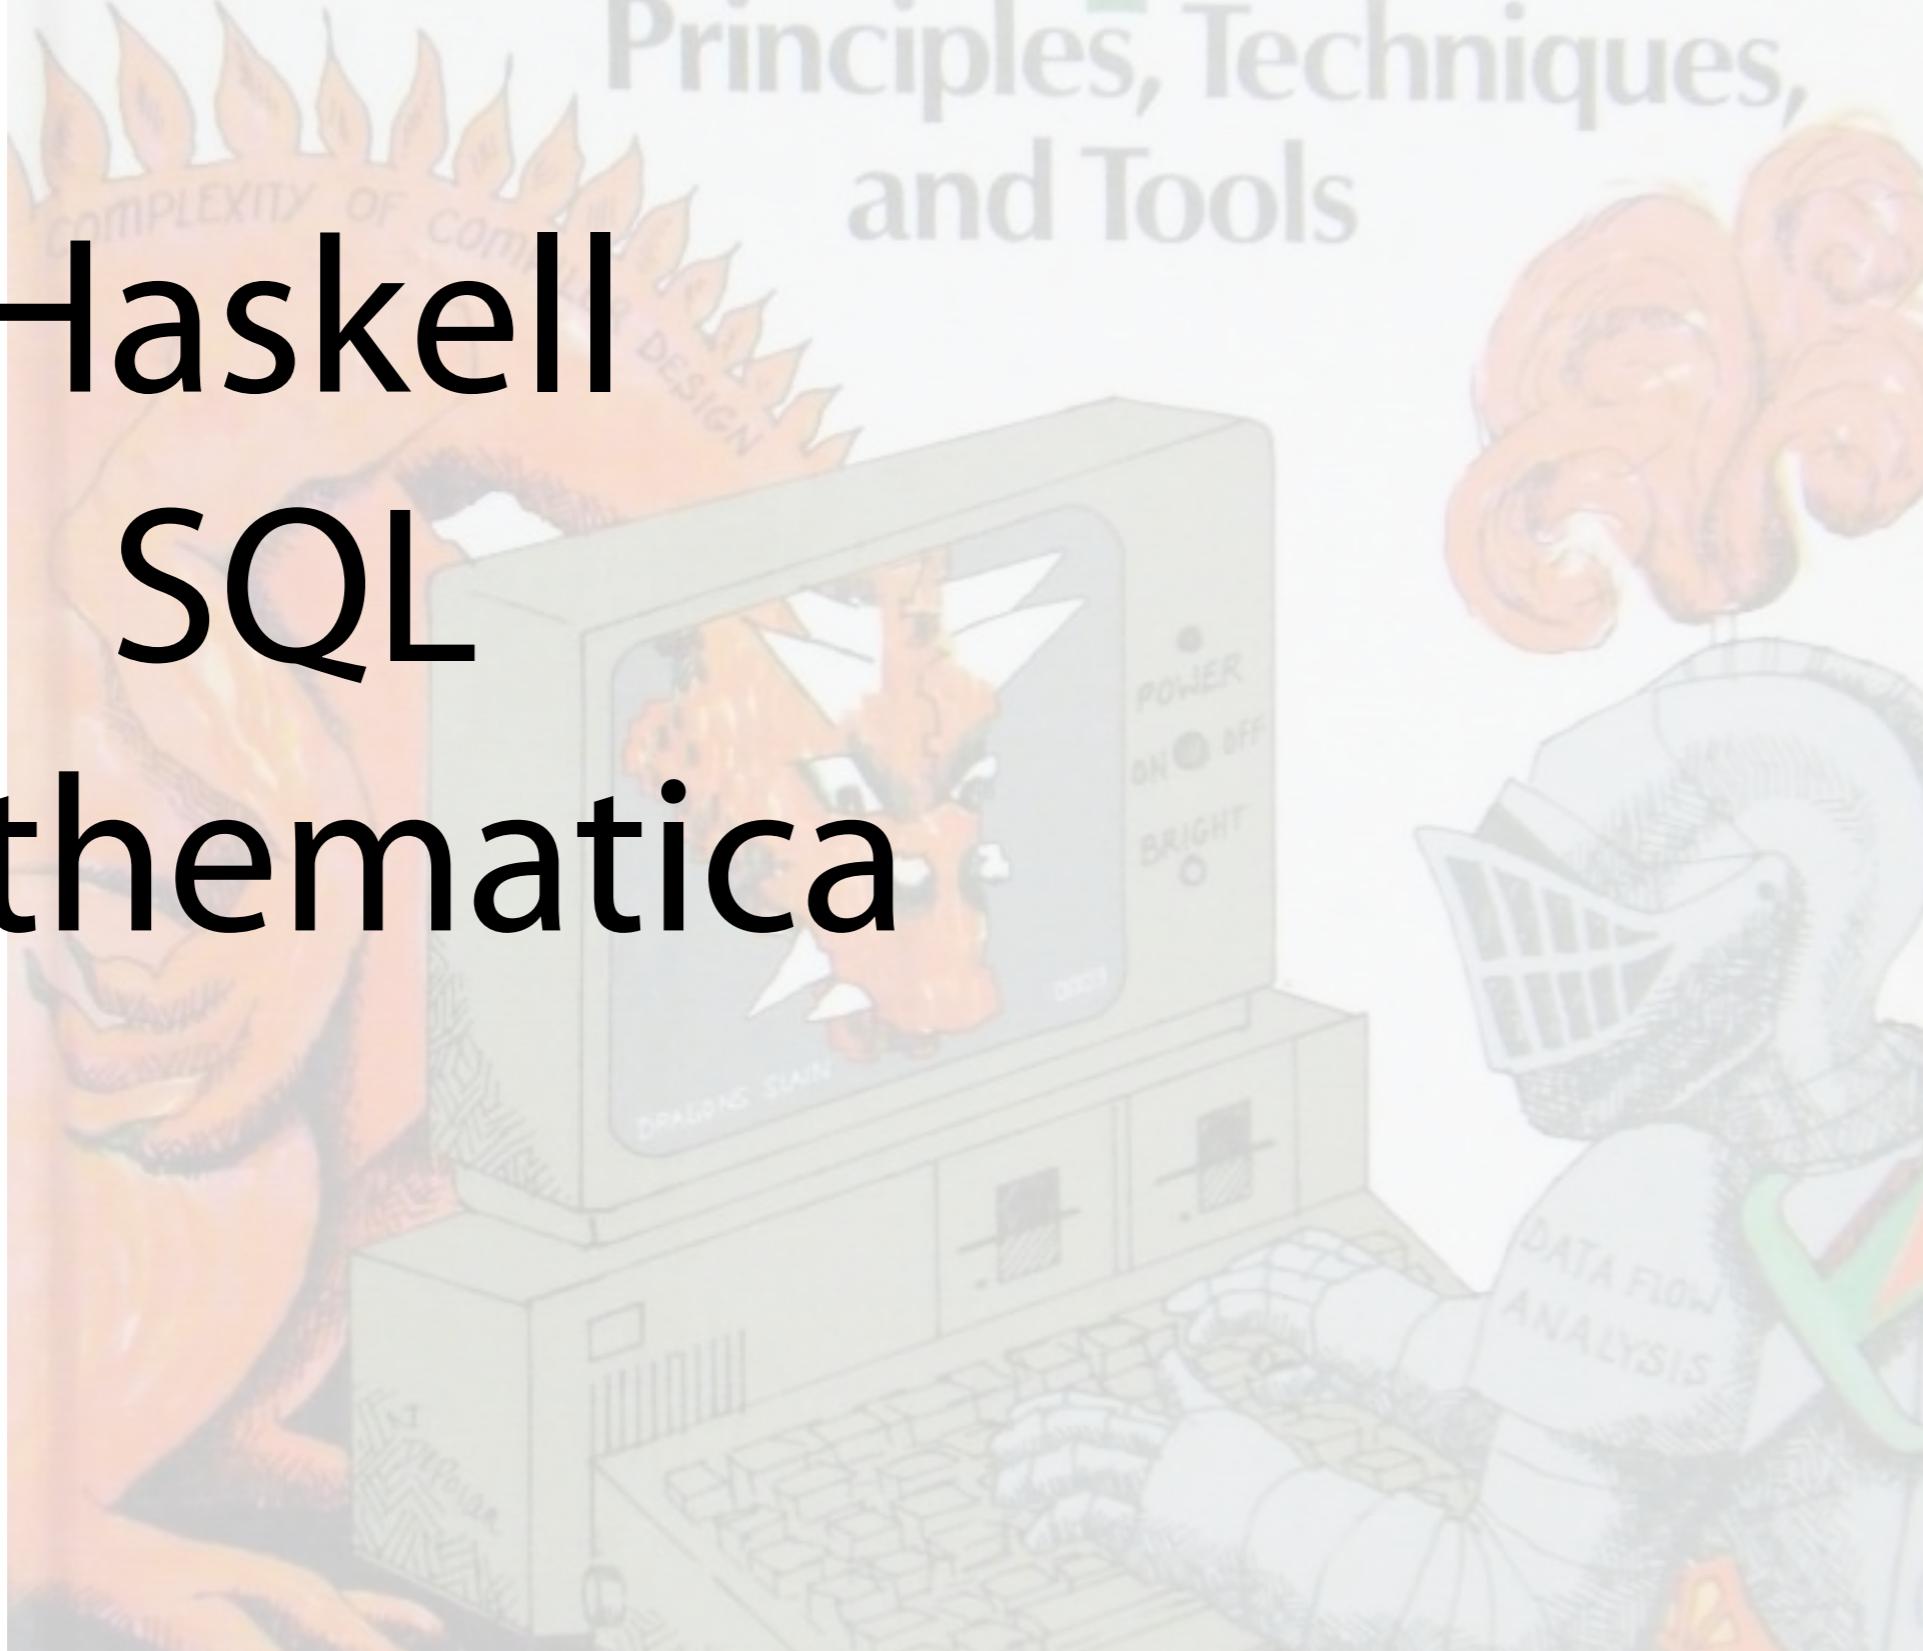


opaque

transparent

Haskell
SQL

Mathematica



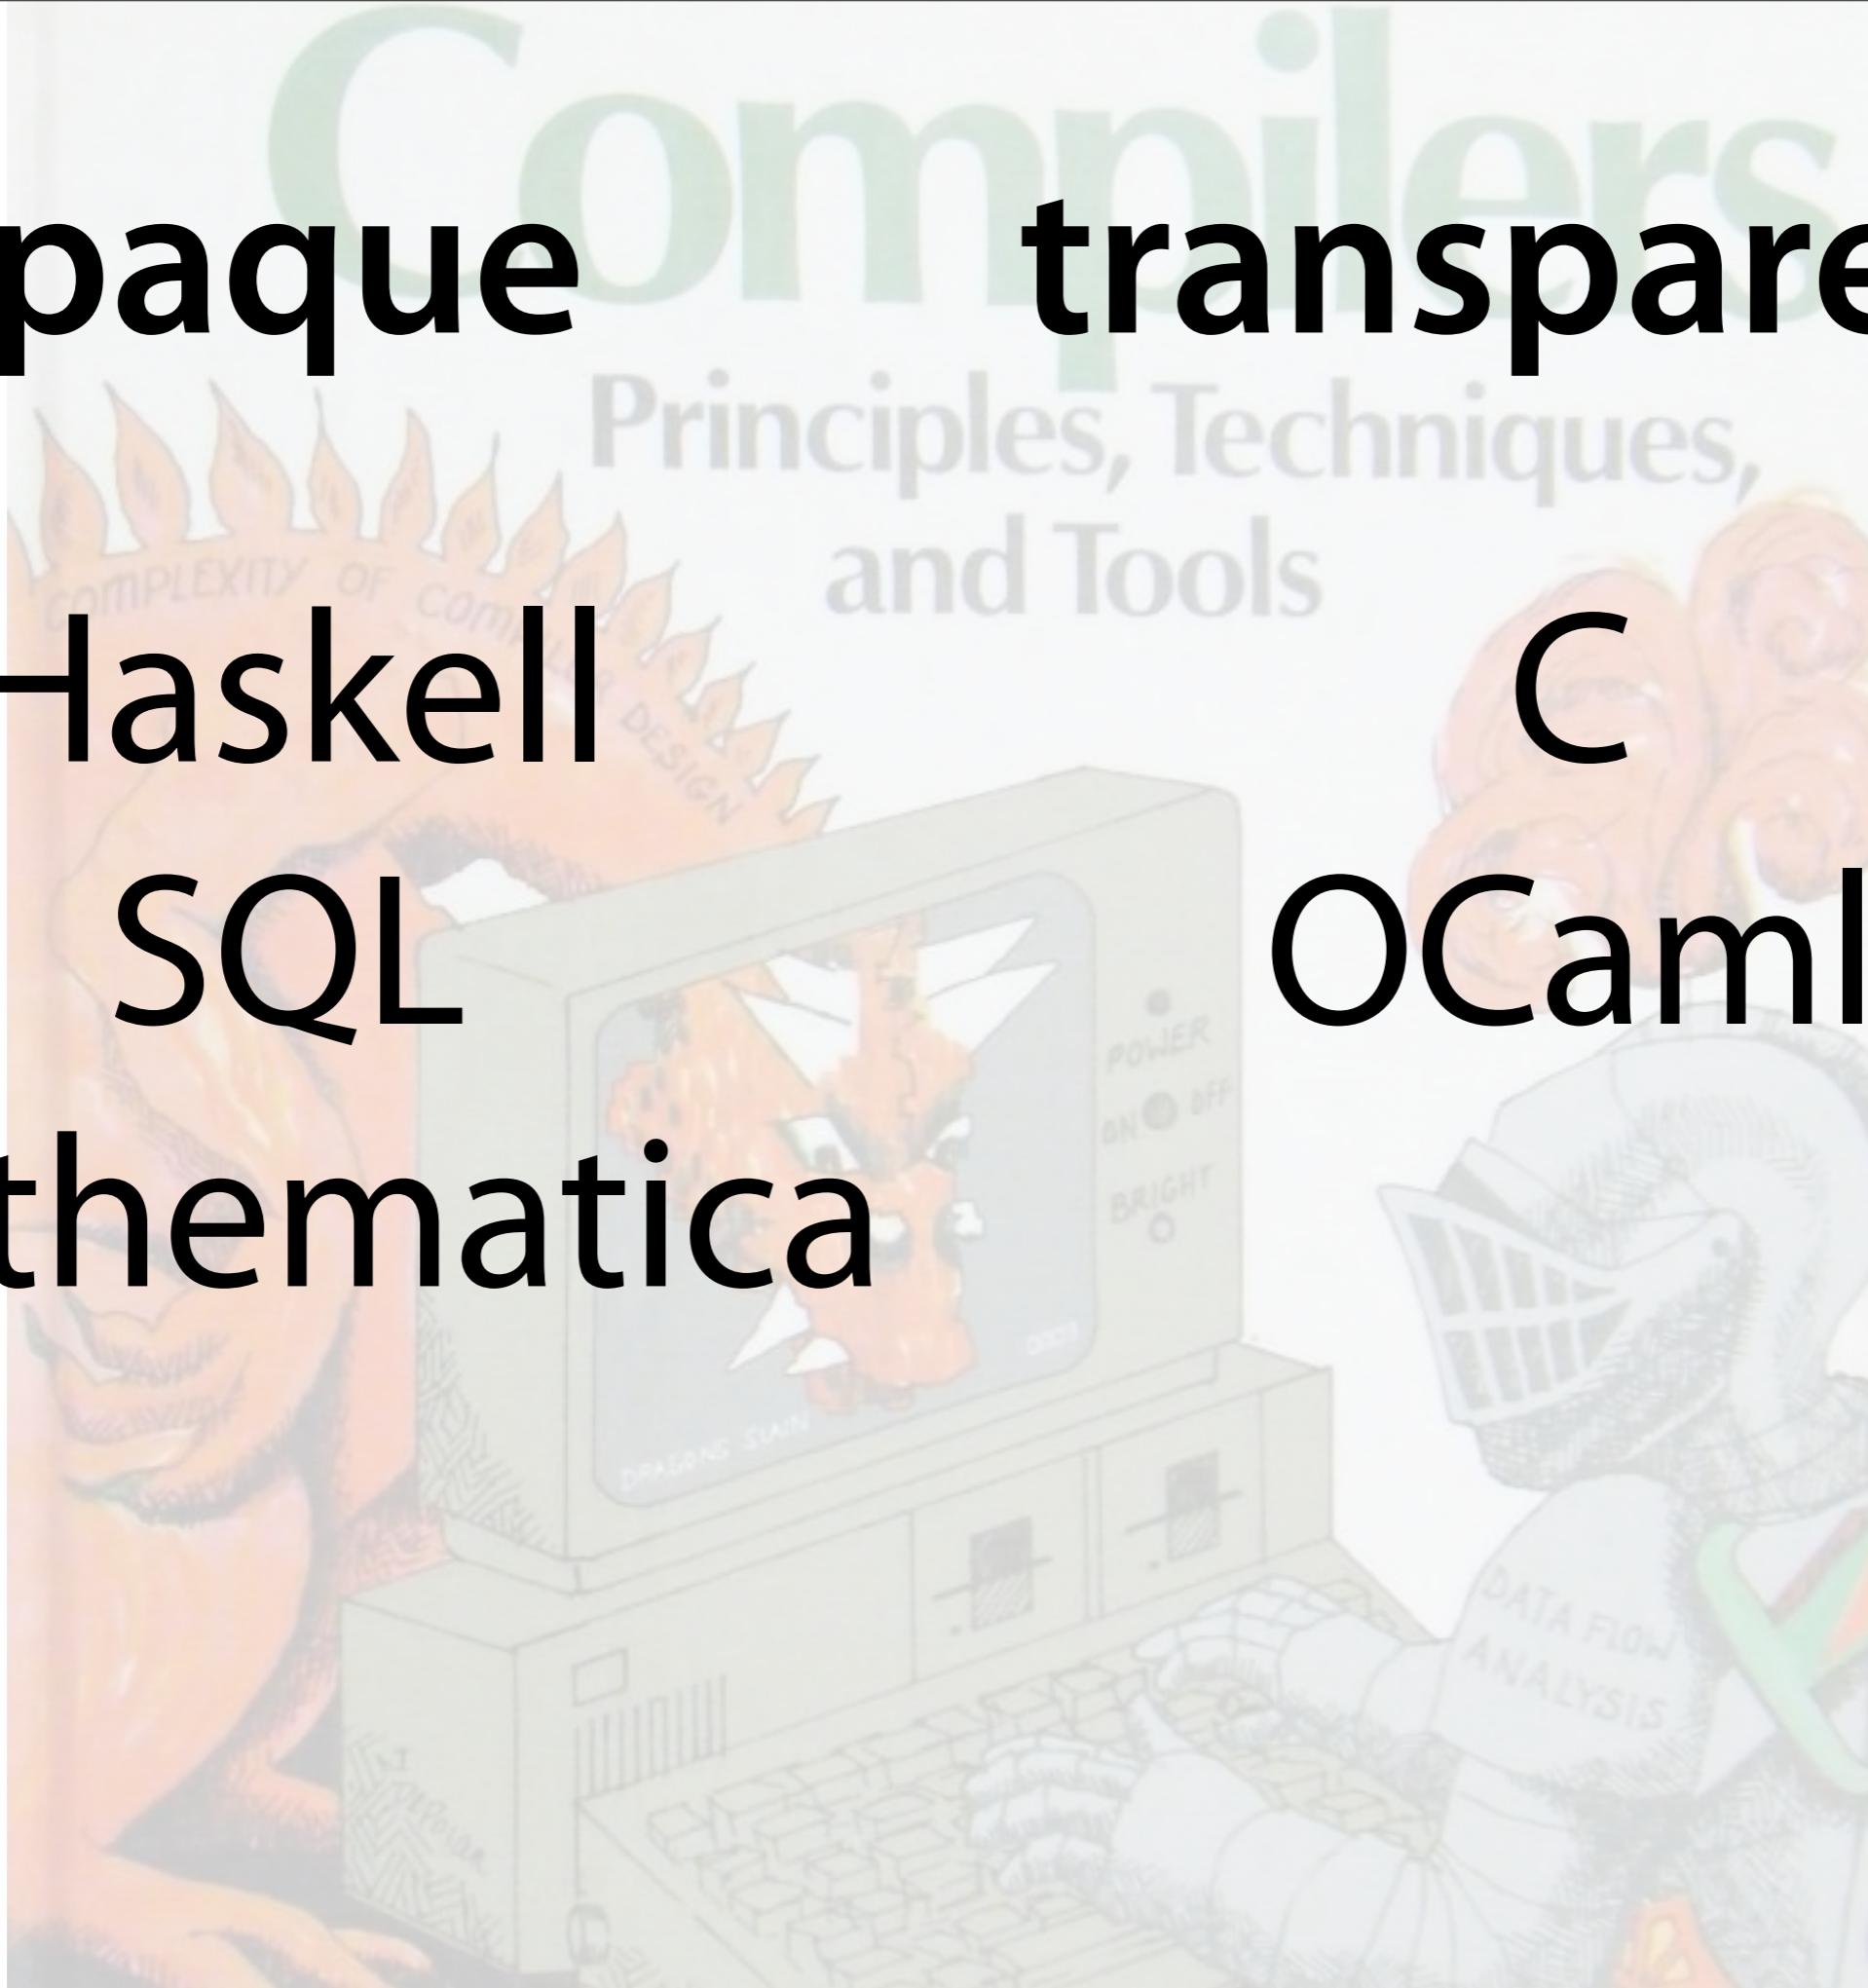
opaque

transparent

Haskell
SQL

Mathematica

C
OCaml



opaque

transparent

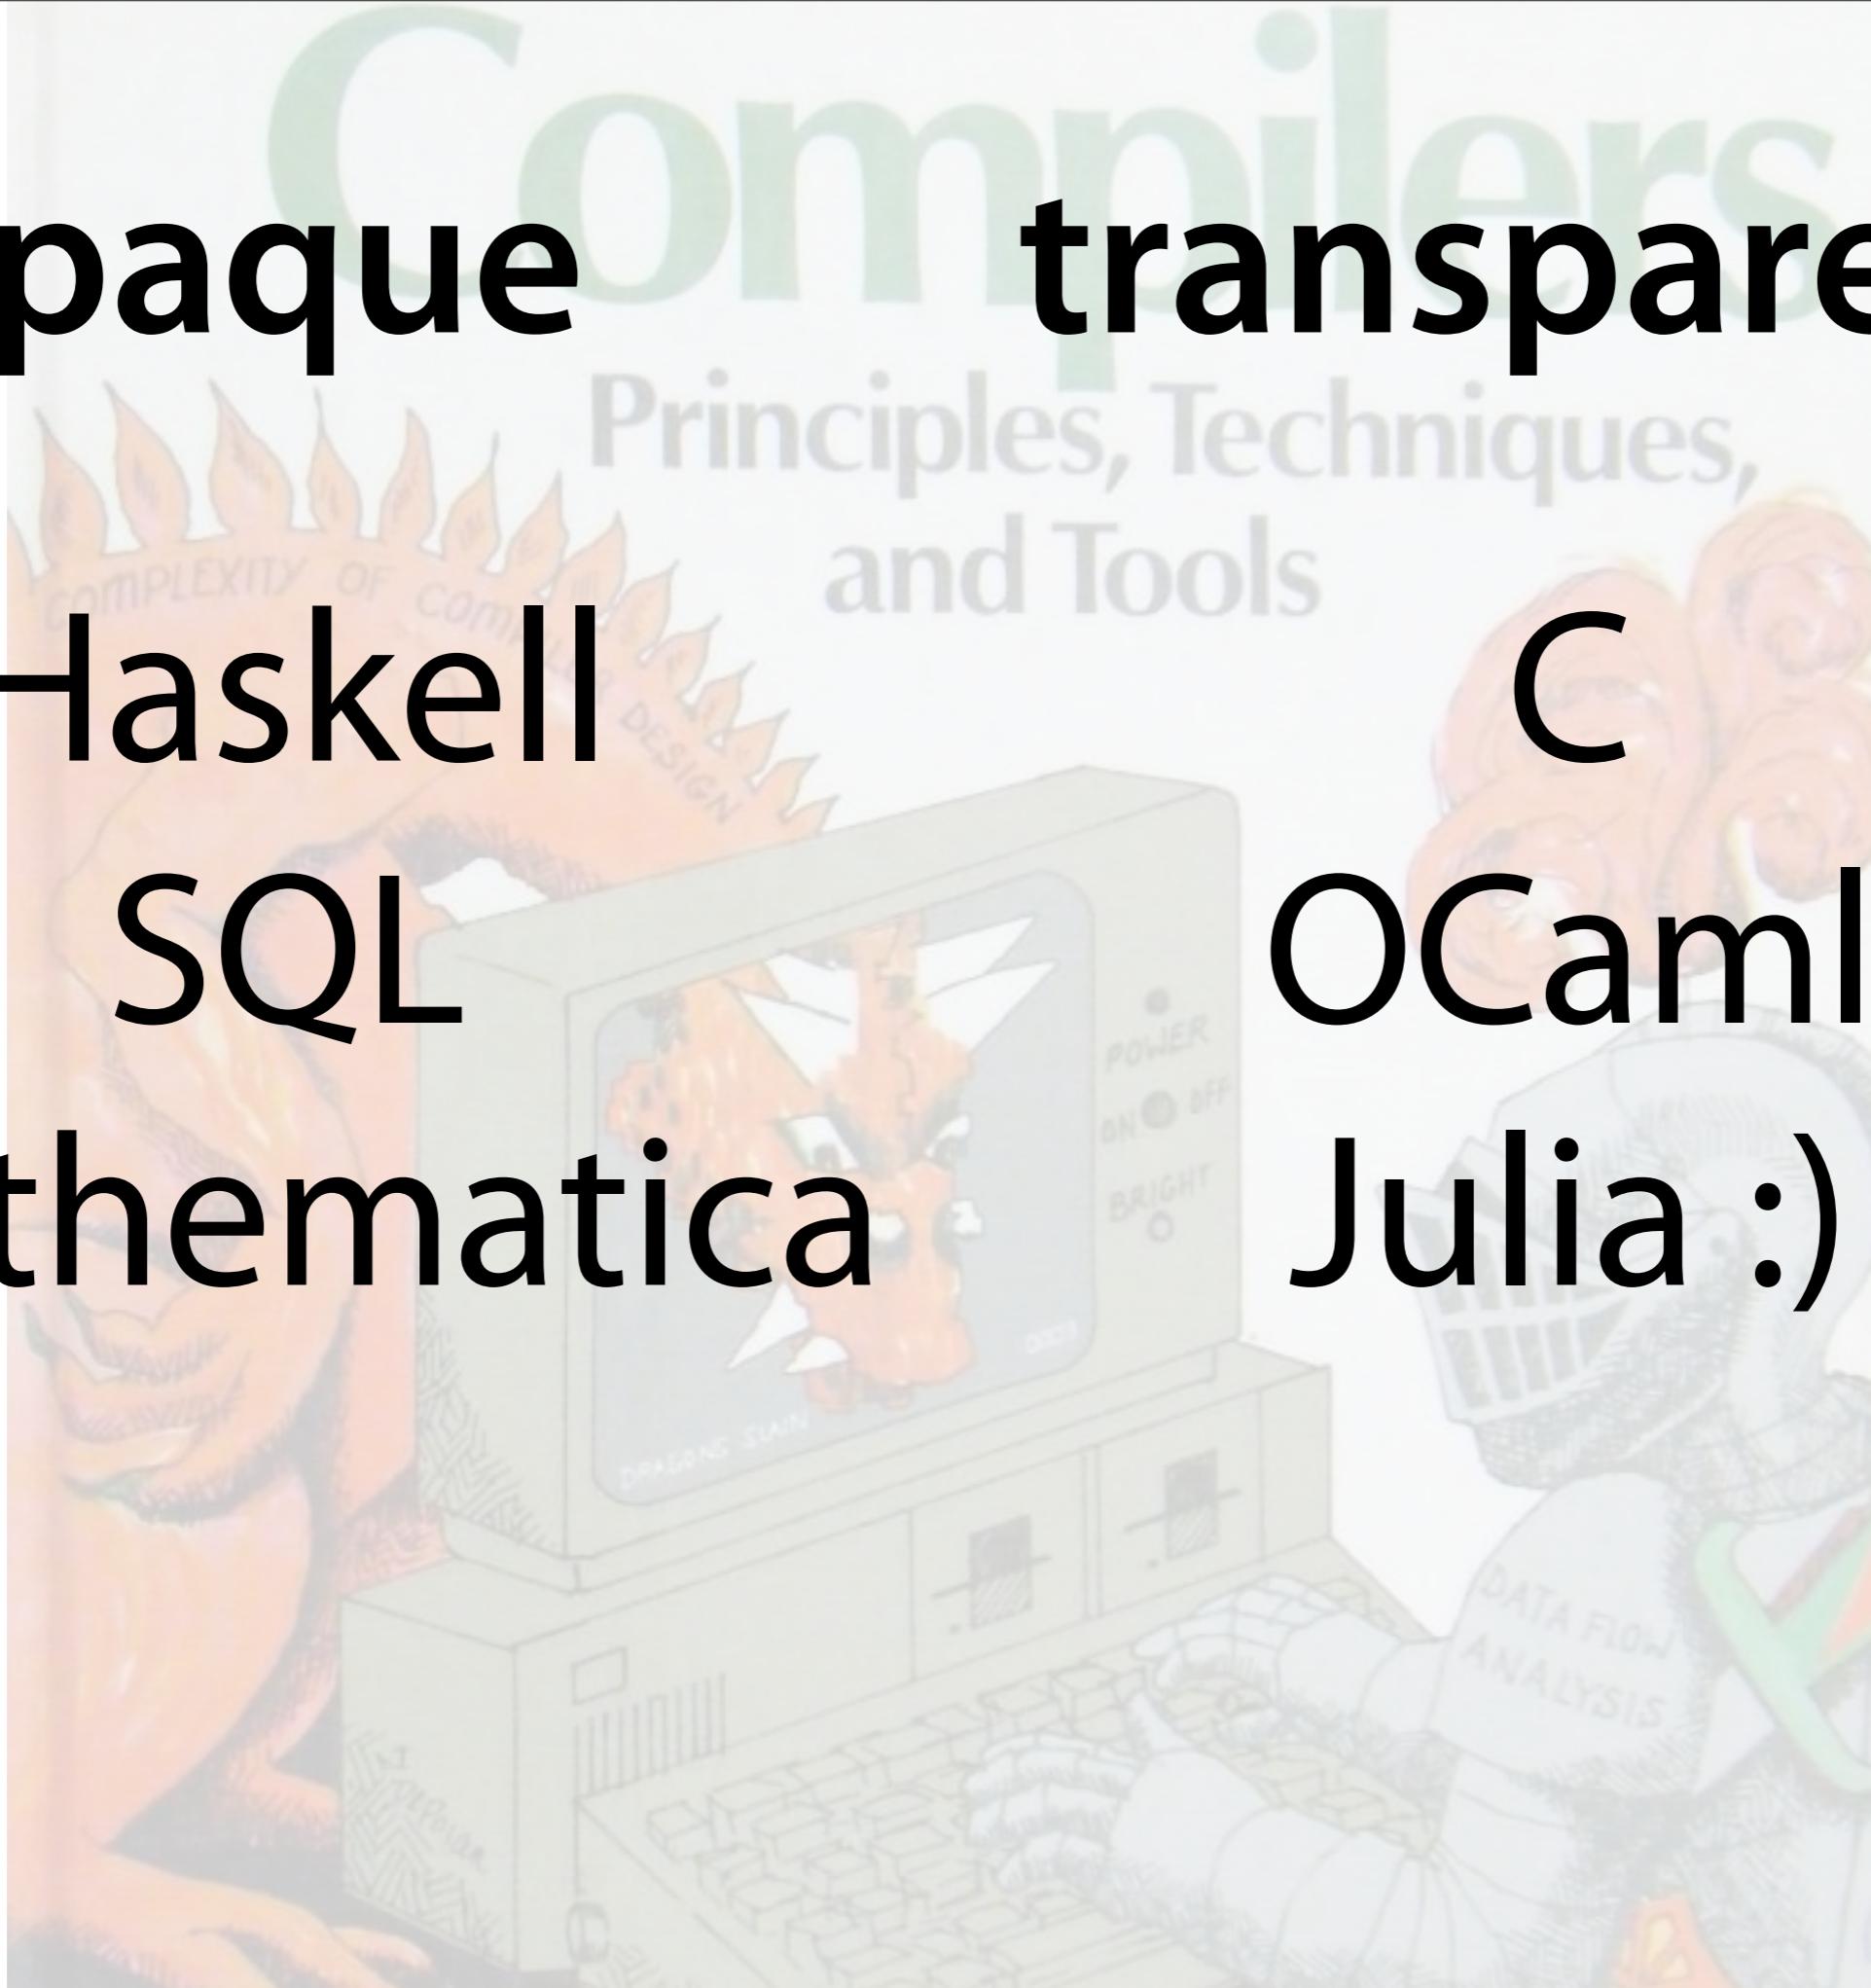
Haskell

SQL

Mathematica

C

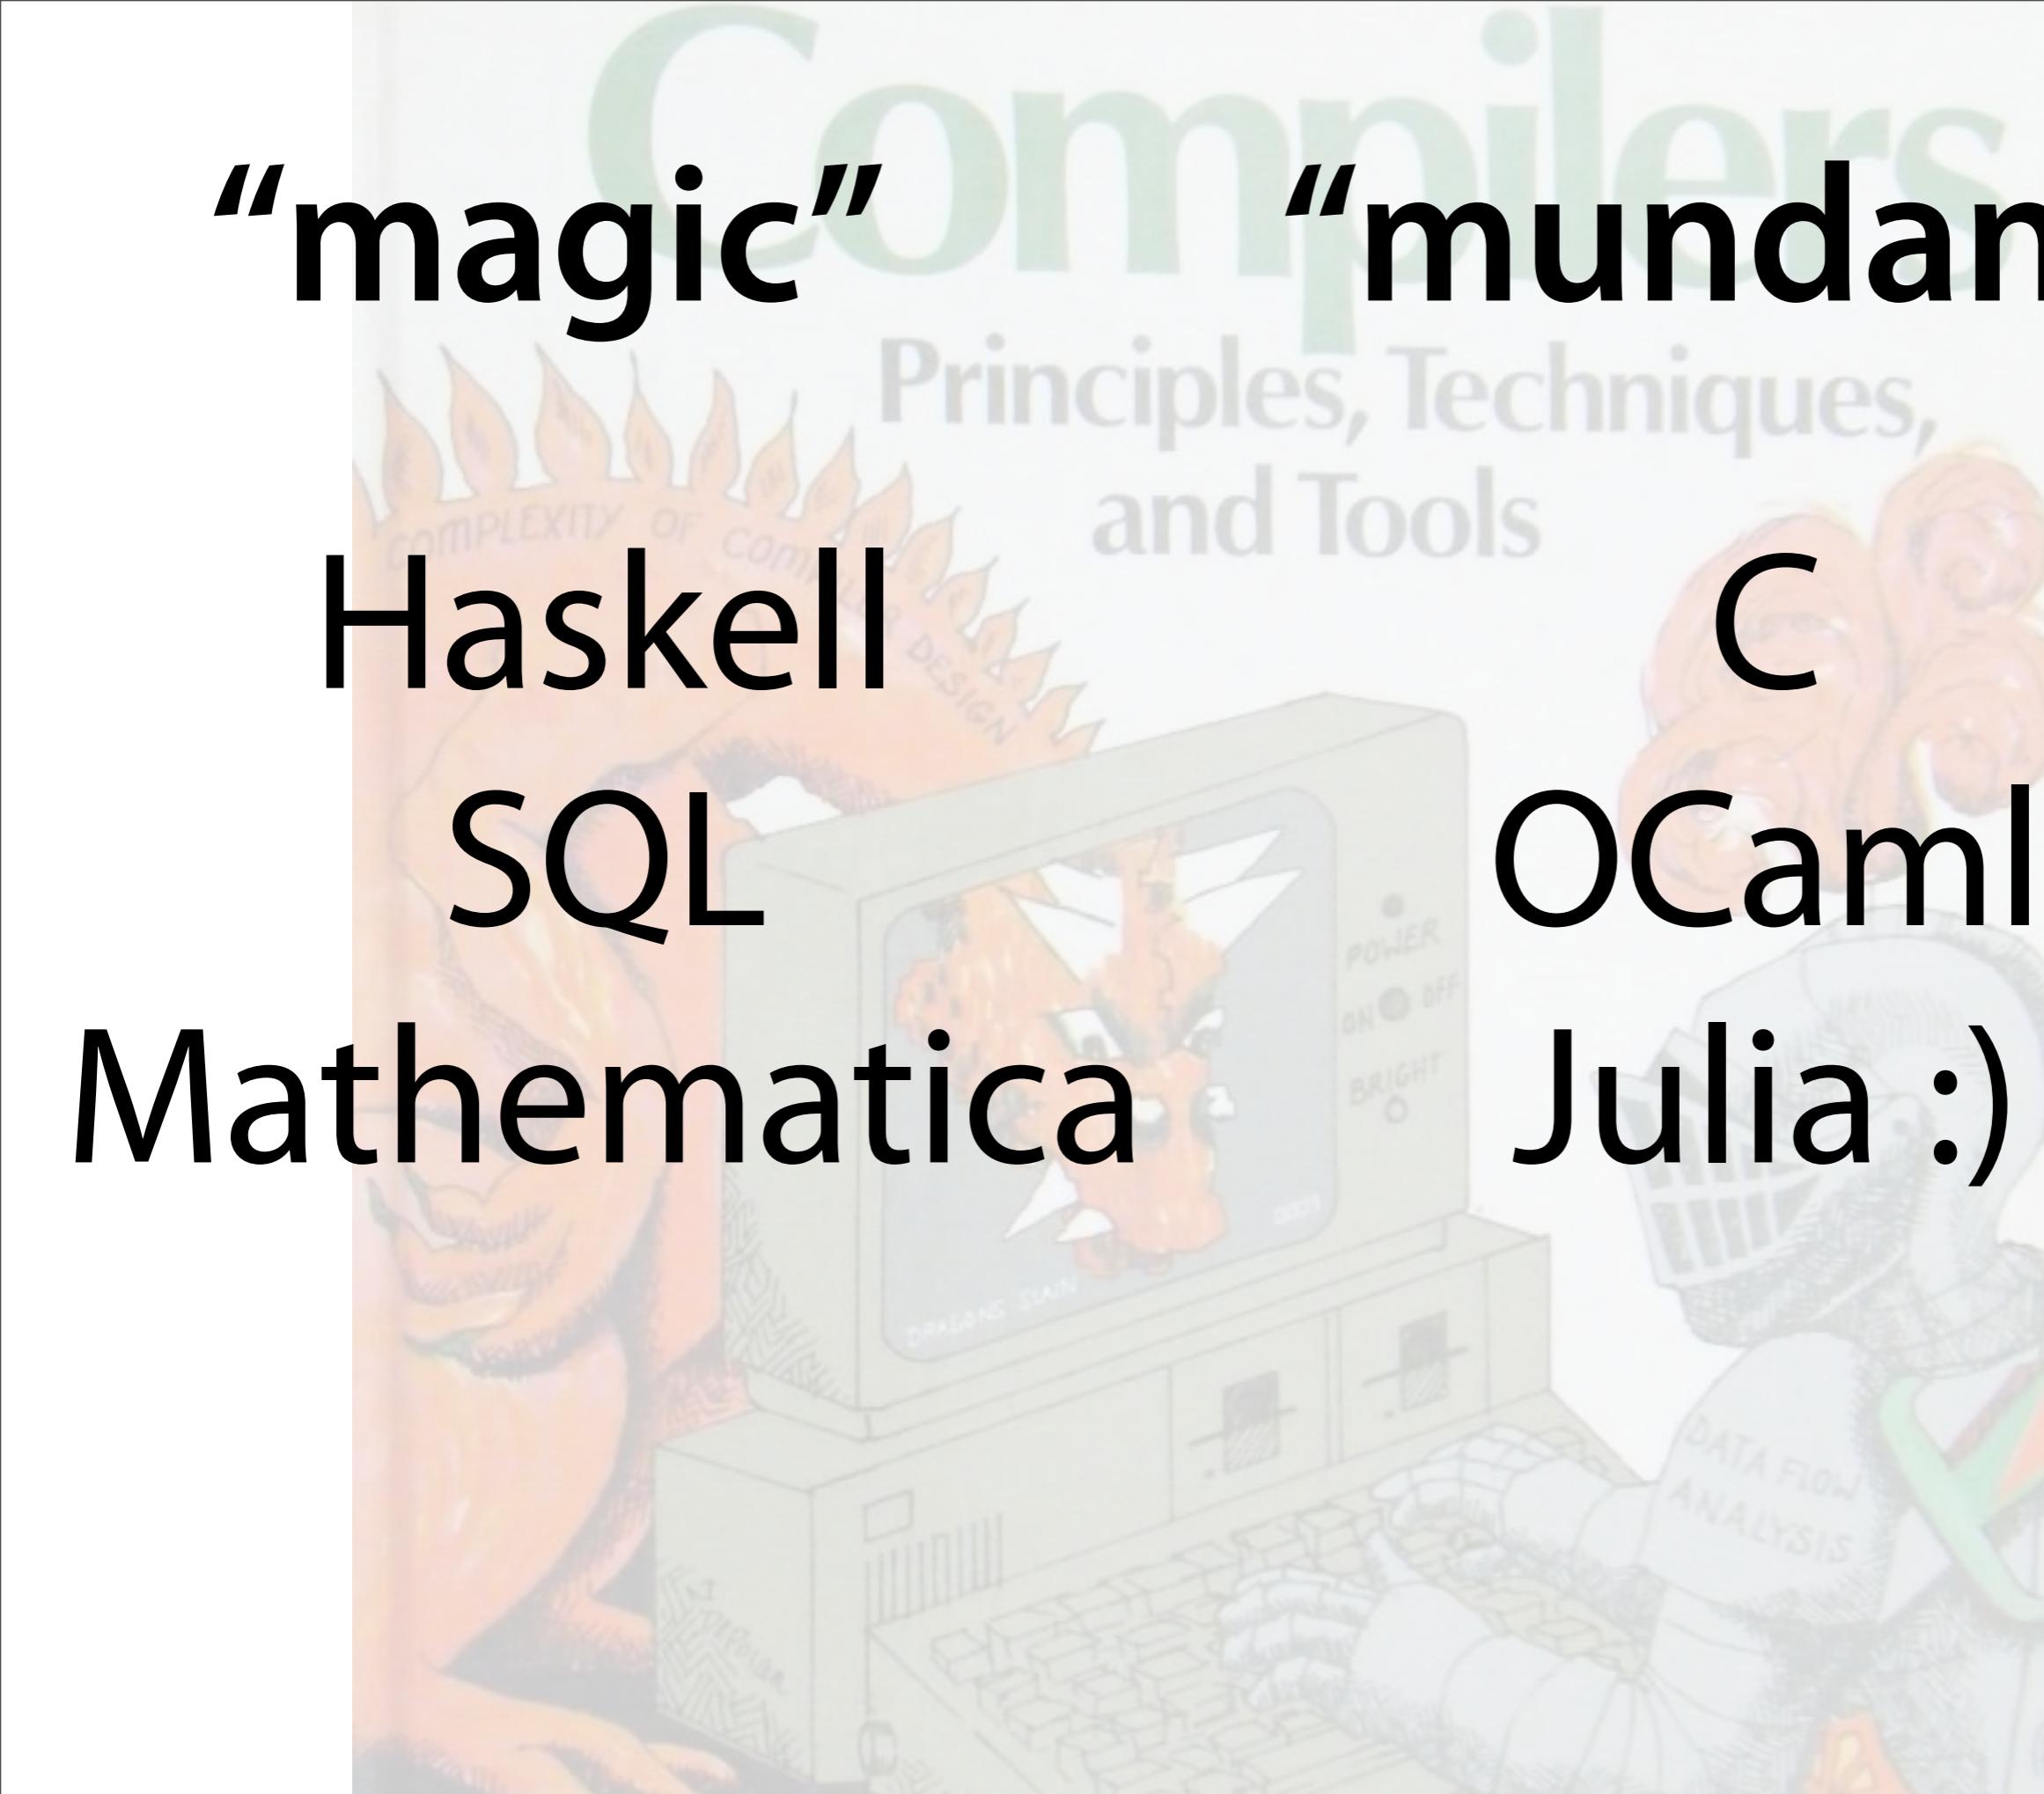
OCaml
Julia :)



“magic” “mundane”

Haskell
SQL
Mathematica

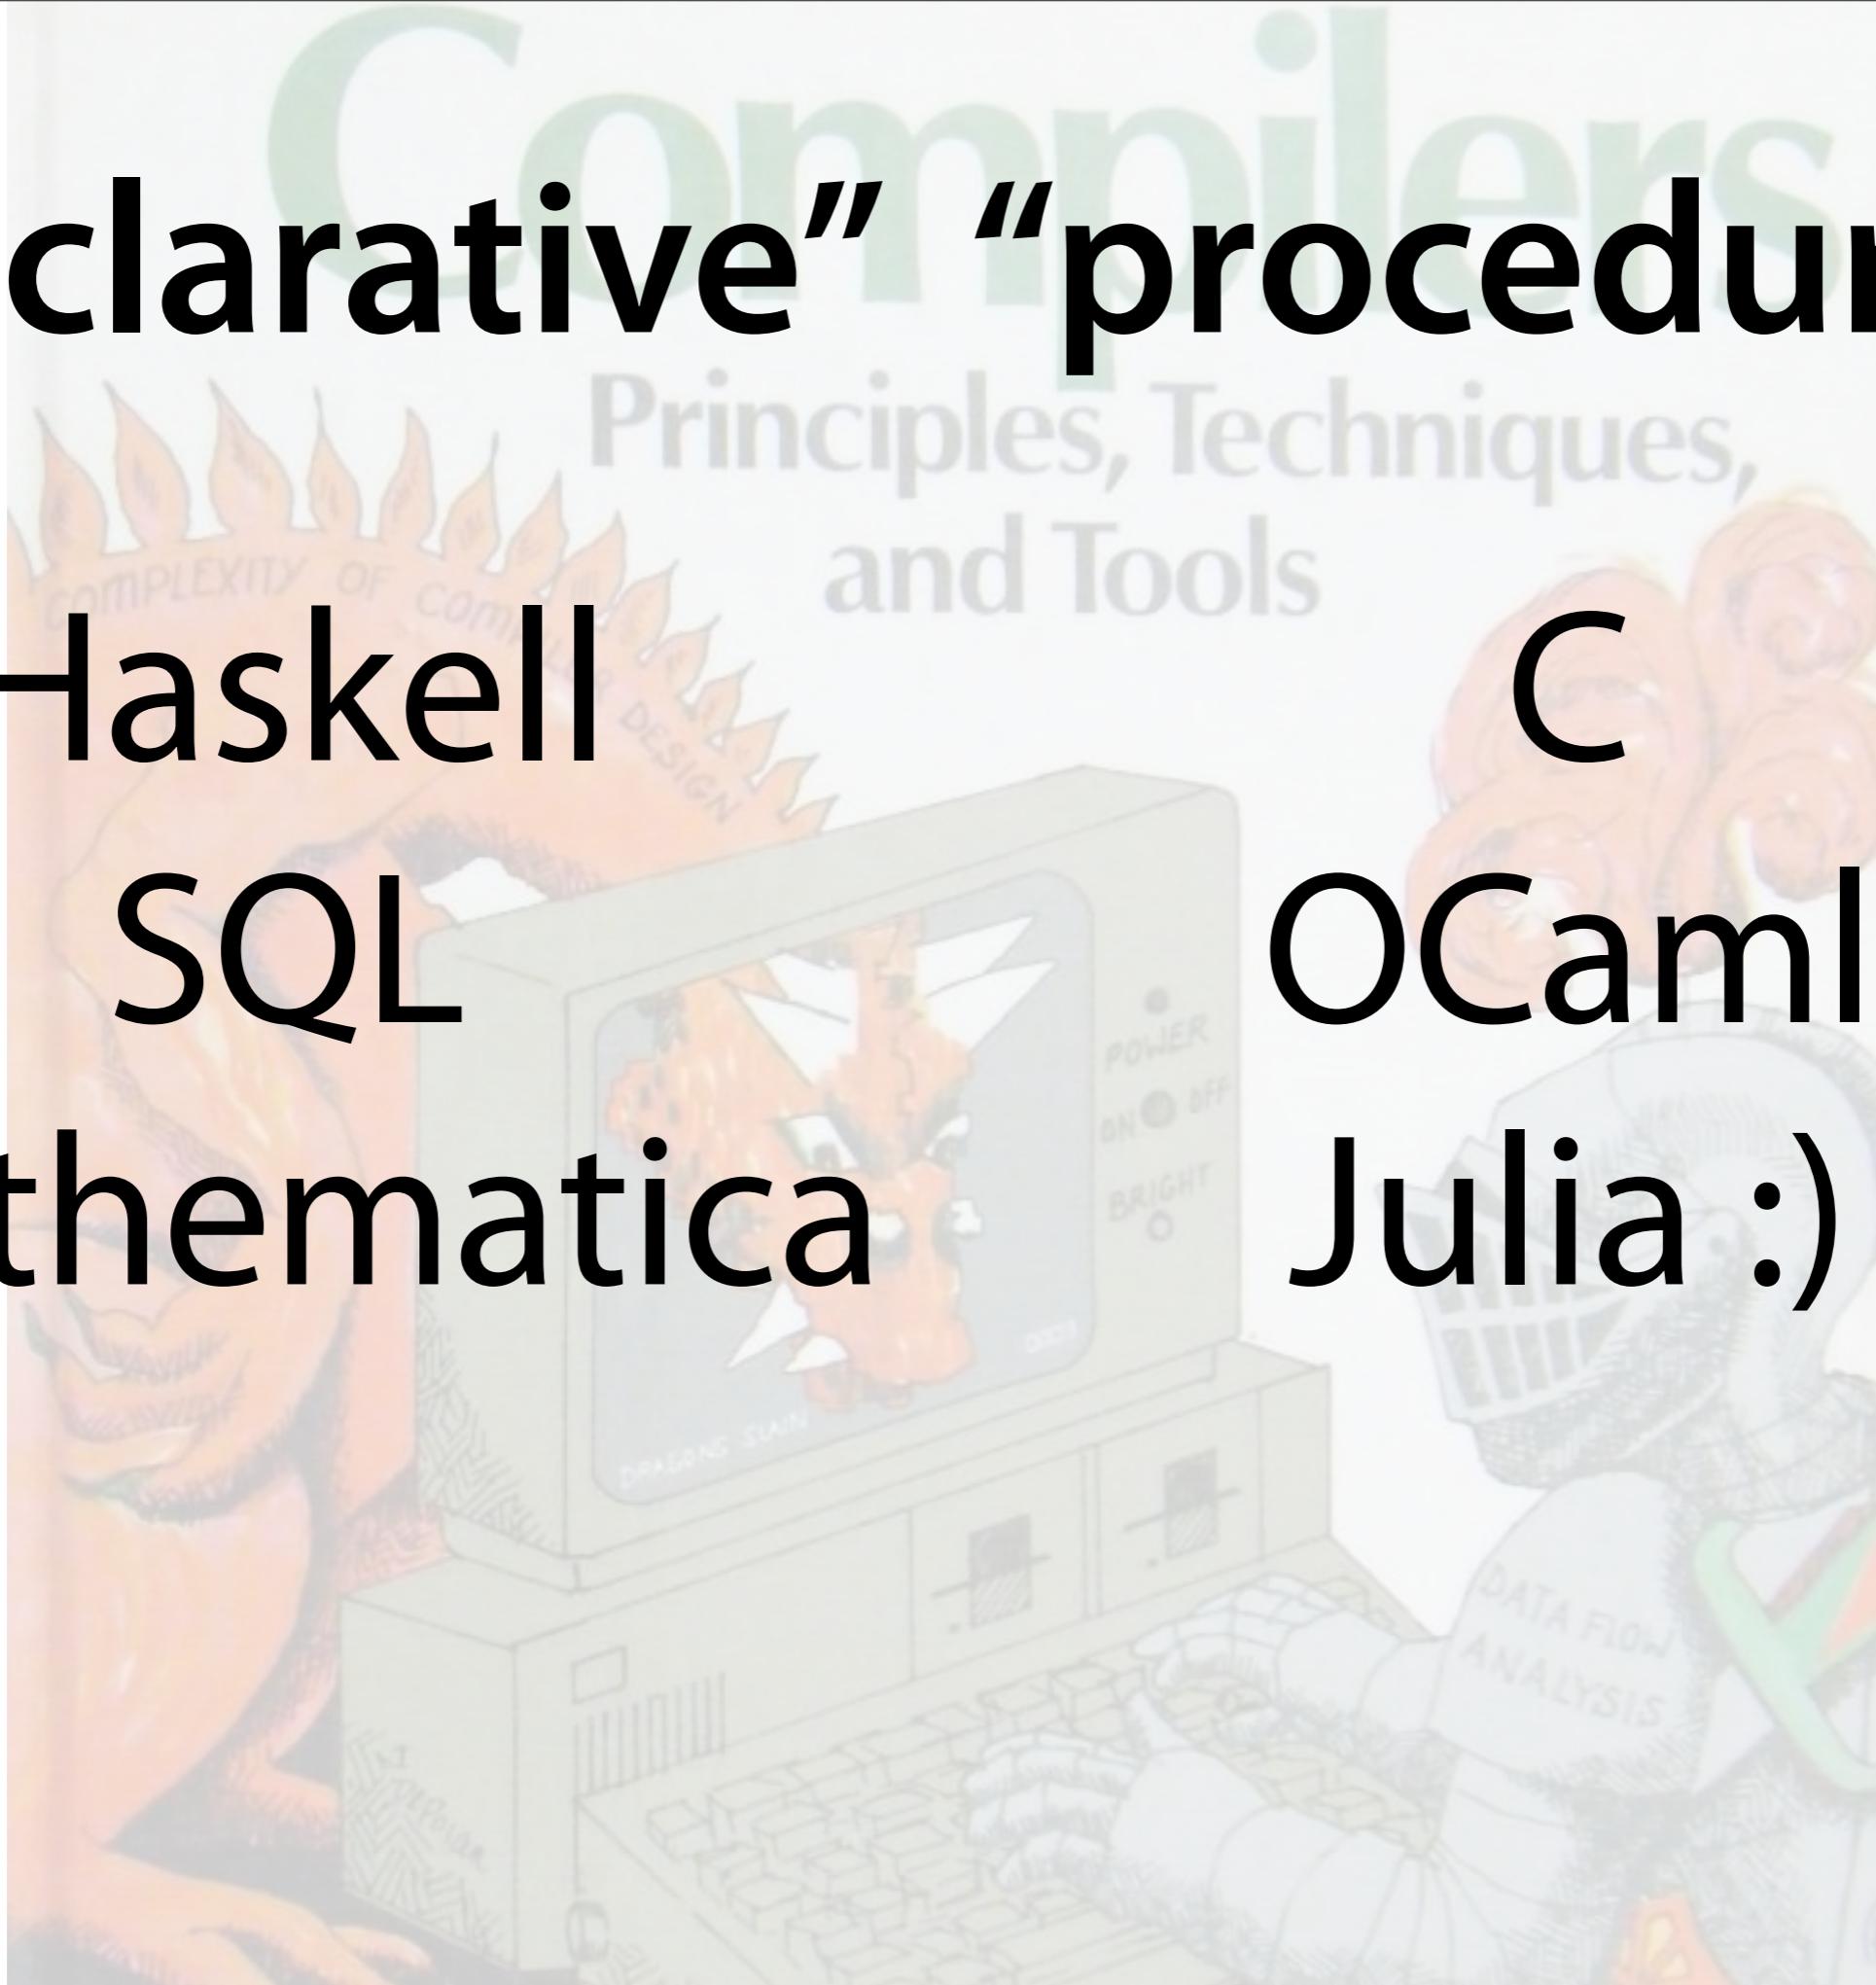
C
OCaml
Julia :)



“declarative” “procedural”

Haskell
SQL
Mathematica

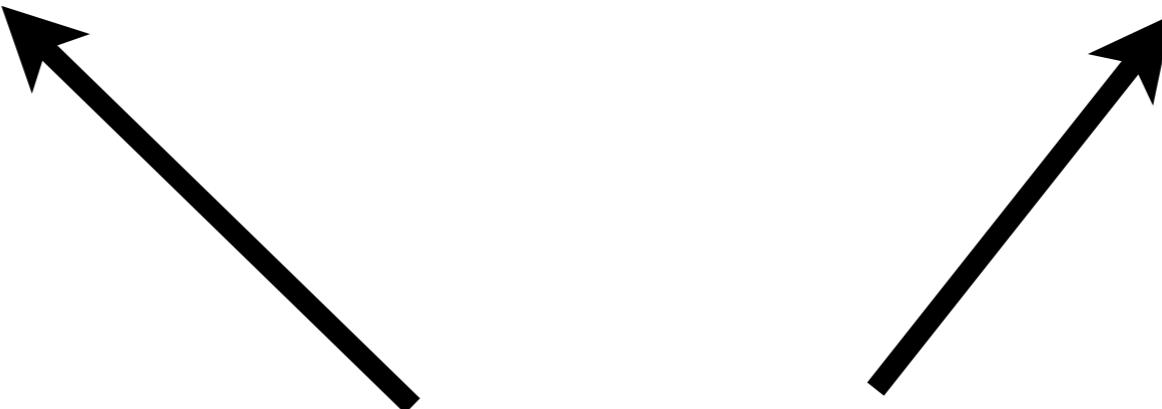
C
OCaml
Julia :)



```
result = A .* B + C .* D + A
```

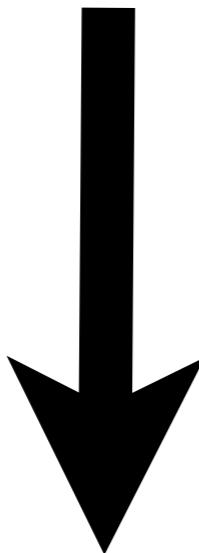
two unnecessary memory allocations

result = A .* B + C .* D + A



A is traversed twice

```
n = length(a)
r = zeros(n)
for i = 1:n
    r[i] = A[i] * B[i] + C[i] * D[i] + A[i]
end
```

$$r = A .* B + C .* D + A$$


“why can’t the compiler
do this for me?”

```
n = length(a)
r = zeros(n)
for i = 1:n
    r[i] = A[i] * B[i] + C[i] * D[i] + A[i]
end
```

simplicity

simplicity

(see https://groups.google.com/d/msg/julia-users/BosGhmT9ei4/ujyLzilA_G4J
for JMW on this)

simplicity

vs.

easiness

simplicity

vs.

easiness

(Cartesian.jl, Devectorize.jl, etc.)

THE
GOOD

THE
BAD

THE GOOD

THE BAD

concise, readable syntax
(for array ops, etc.)

THE GOOD

concise, readable syntax
(for array ops, etc.)

THE BAD

start from
bad code and optimize

THE GOOD

concise, readable syntax
(for array ops, etc.)

start from
bad code and optimize

THE BAD

optimization can
often be significant work

THE GOOD

concise, readable syntax
(for array ops, etc.)

start from
bad code and optimize

compiler is **simple**

THE BAD

optimization can
often be significant work

THE GOOD

concise, readable syntax
(for array ops, etc.)

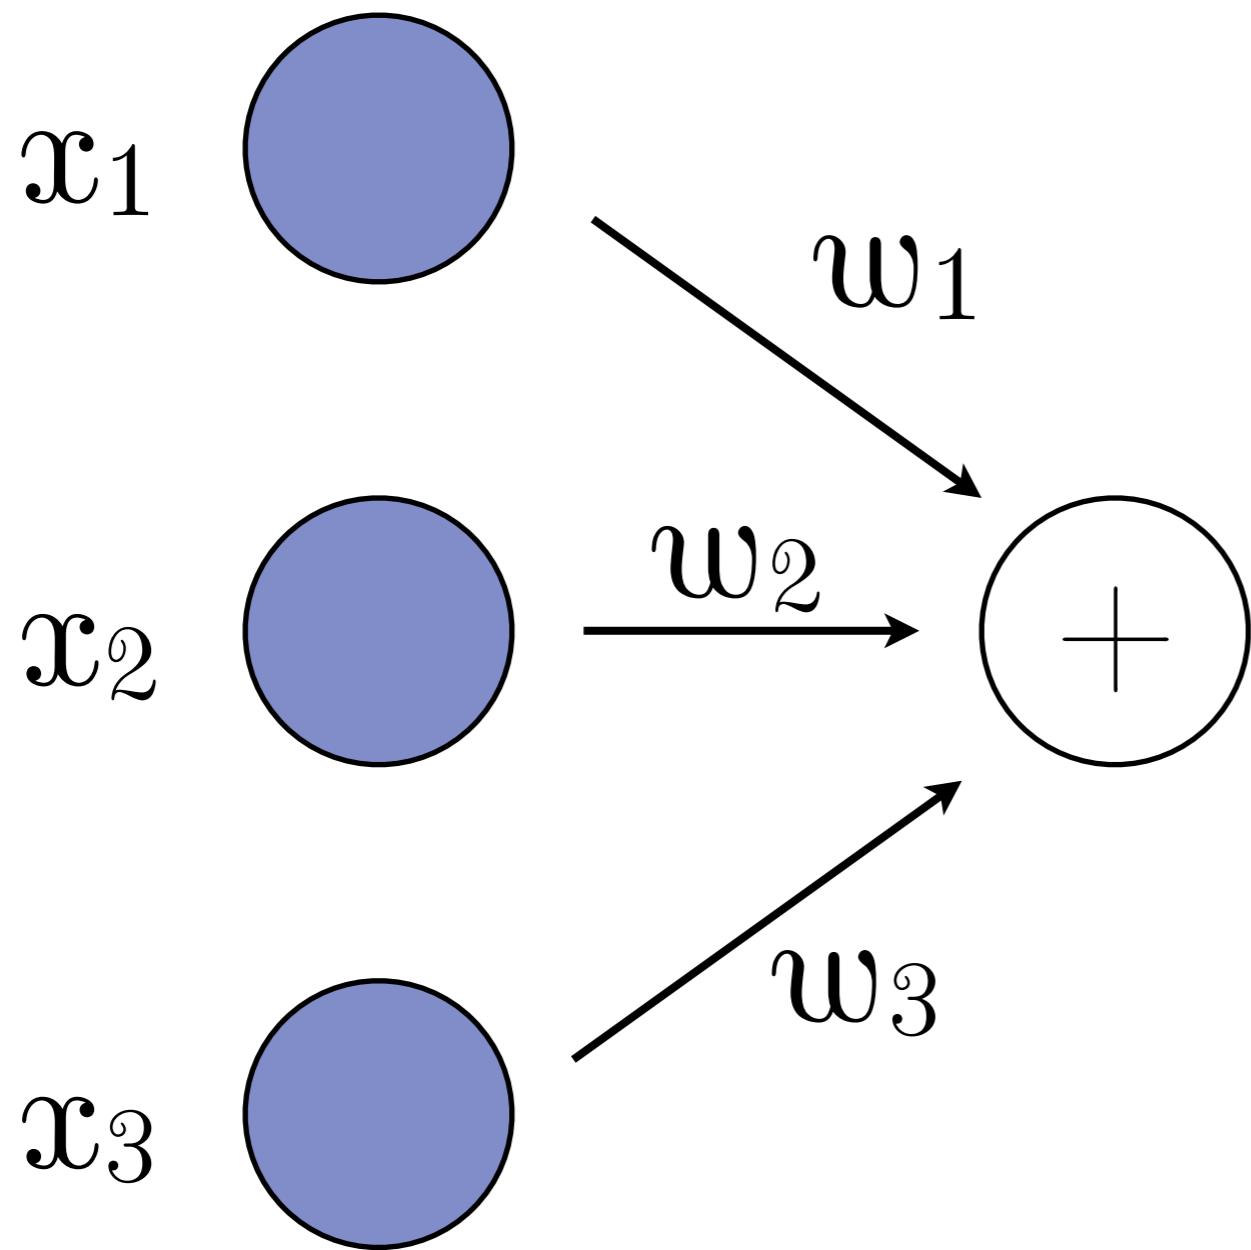
start from
bad code and optimize

compiler is **simple**

THE BAD

optimization can
often be significant work

sometimes getting
performance is not **easy**



```
if result > b  
    classify 1  
else  
    classify -1  
end
```

```
yhat = dot(w,data[:,i]) >= 0 ? 1 : -1
if yhat == -1 && labels[i] == 1
    w = w + data[:,i]
    nmistakes += 1
elseif yhat == 1 && labels[i] == -1
    w = w - data[:,i]
    nmistakes += 1
end
```

```
yhat = dot(w,data[:,i]) >= 0 ? 1 : -1
if yhat == -1 && labels[i] == 1
    w = w + data[:,i]
    nmistakes += 1
elseif yhat == 1 && labels[i] == -1
    w = w - data[:,i]
    nmistakes += 1
end
```

```
yhat = dot(w,data[:,i]) >= 0 ? 1 : -1
if yhat == -1 && labels[i] == 1
    w = w + data[:,i]
    nmistakes += 1
elseif yhat == 1 && labels[i] == -1
    w = w - data[:,i]
    nmistakes += 1
end
```

each of these is a copy!

```
yhat = dot(w,slice(data,:,i)) >= 0
    ? 1 : -1
if yhat == -1 && labels[i] == 1
    w = w + slice(data,:,i)
    nmistakes += 1
elseif yhat == 1 && labels[i] == -1
    w = w - slice(data,:,i)
    nmistakes += 1
end
```

```
yhat = dot(w,slice(data,:,i)) >= 0
    ? 1 : -1
if yhat == -1 && labels[i] == 1
    w = w + slice(data,:,i)
    nmistakes += 1
elseif yhat == 1 && labels[i] == -1
    w = w - slice(data,:,i)
    nmistakes += 1
end
```

still not fast :(



lindahua / ArrayViews.jl

```
# get initial results
res = pmap( $\sigma$  -> crossvalidate(data, labels,  $\sigma$ ),
 $\sigma$ vals)
```

```
# get initial results
res = pmap( $\sigma$  -> crossvalidate(data, labels,  $\sigma$ ),
           sigma)
```

```
# get initial results
res = pmap(σ -> crossvalidate(data, labels, σ),
σvals)
```

slow :(

```
# get initial results
res = pmap( $\sigma$  -> crossvalidate(data, labels,  $\sigma$ ),
 $\sigma$ vals)
```

```
# get initial results
res = pmap(σ -> crossvalidate(data, labels, σ),
σvals)
```

UTF8! :)

a ≠ b

a ≠ b

a ≤ b

a ≠ b

a ≤ b

a · b

a ≠ b

a ≤ b

a · b

a × b

a ≠ b

a ≤ b

a · b

a × b

a ∪ b

$$(\text{ } \circ_{\square} \text{ } \circ) \text{ } \prime \sim \overline{\perp \perp} = 3$$

```
# get initial results
res = pmap( $\sigma \rightarrow$  crossvalidate(data, labels,  $\sigma$ ),
            $\sigma$ vals)
```

```
include("mycode.jl")  
  
data = load_data("data_source")  
  
map(do_lots_of_work, data)
```

```
addprocs(4)
```

```
@everywhere include("mycode.jl")
```

```
data = load_data("data_source")
```

```
pmap(do_lots_of_work, data)
```

```
trials = {}
for filename in filenames
    data = readdlm(filename, Int)
    for mu in mus
        noised = noise(data, mu)
        for theta in thetas
            time, denoised = denoise(noised, theta)
            trial = {"file" => filename,
                      "θ" => theta,
                      "μ" => mu,
                      "time" => time}
            push!(trials, trial)
    end
end
end
```

```
trials = @parallel vcat for filename in filenames
    data = readdlm(filename, Int)
    @parallel vcat for mu in mus
        noised = noise(data, mu)
        @parallel vcat for theta in thetas
            time, denoised = denoise(noised, theta)
            trial = {"file" => filename,
                      "θ" => theta,
                      "μ" => mu,
                      "time" => time}
        trial
    end
end
end
```

THE
GOOD

THE
BAD

lambdas

THE
GOOD

THE
BAD

lambdas

unicode identifiers

THE
GOOD

THE
BAD

lambdas

unicode identifiers

simple and easy
parallel programming

THE GOOD

THE BAD

lambdas

... that aren't
optimized

unicode identifiers

simple **and** easy
parallel programming

THE GOOD

THE BAD

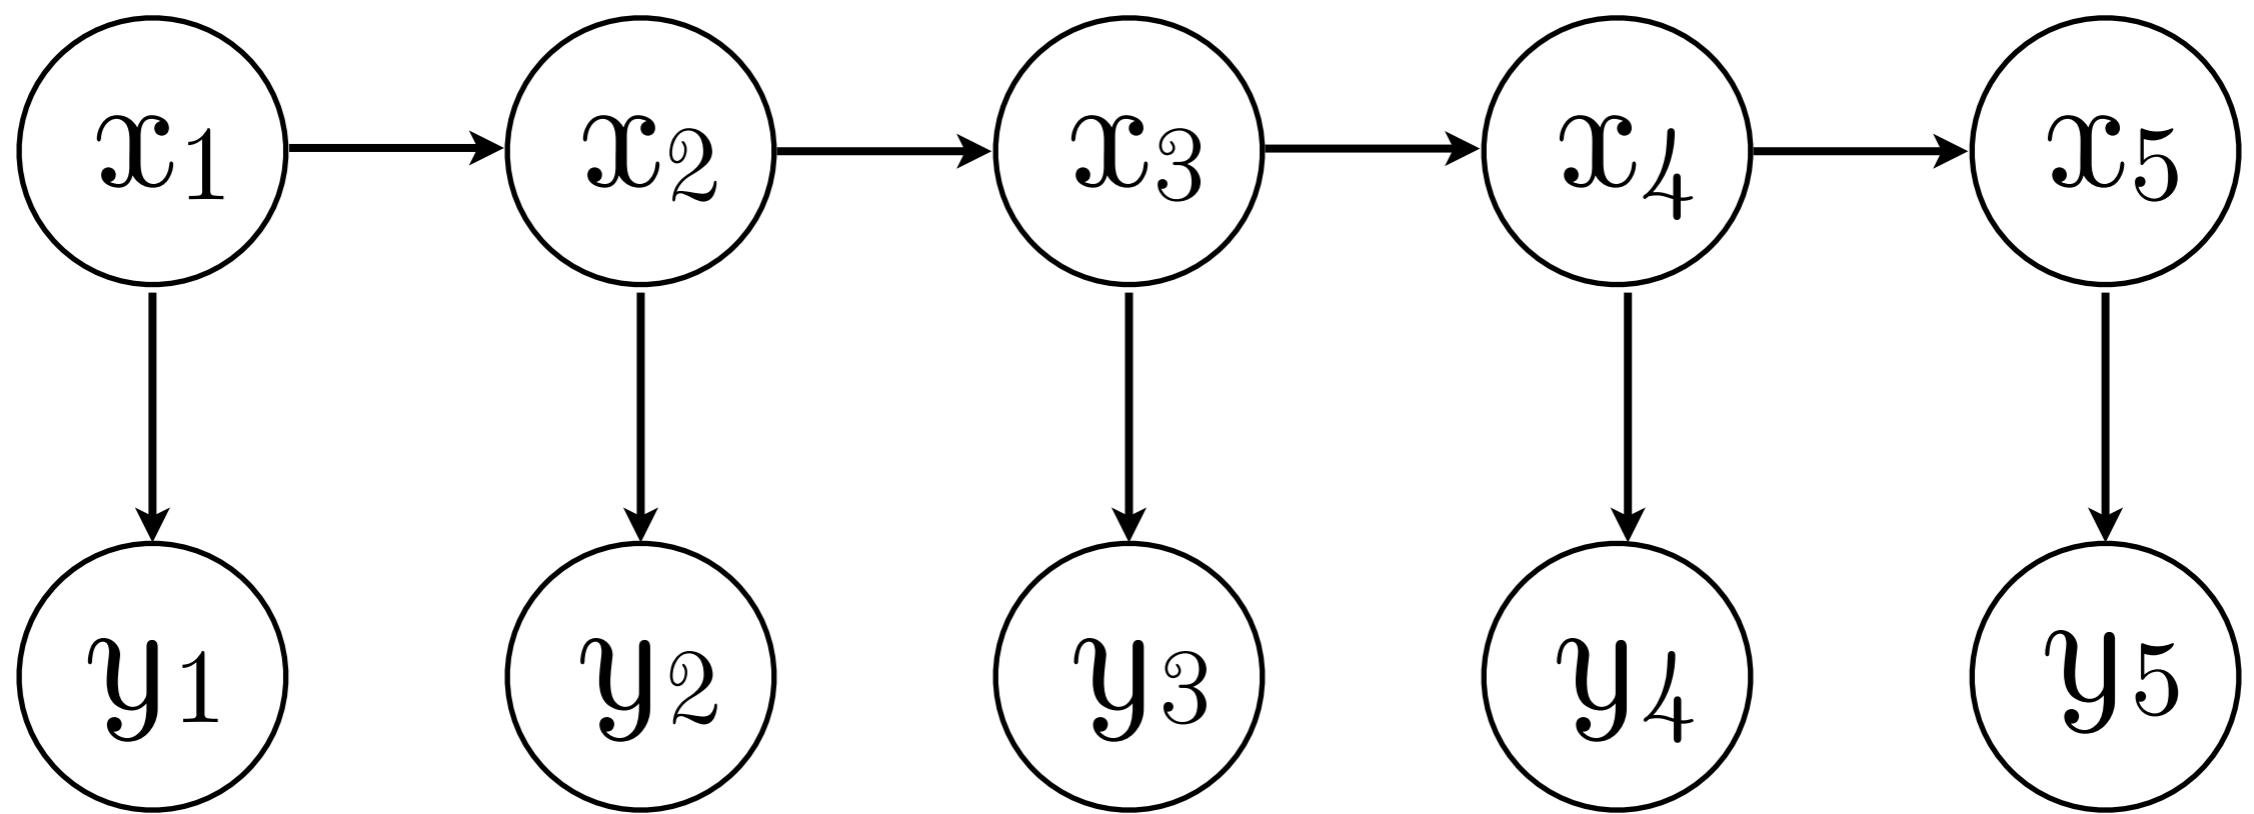
lambdas

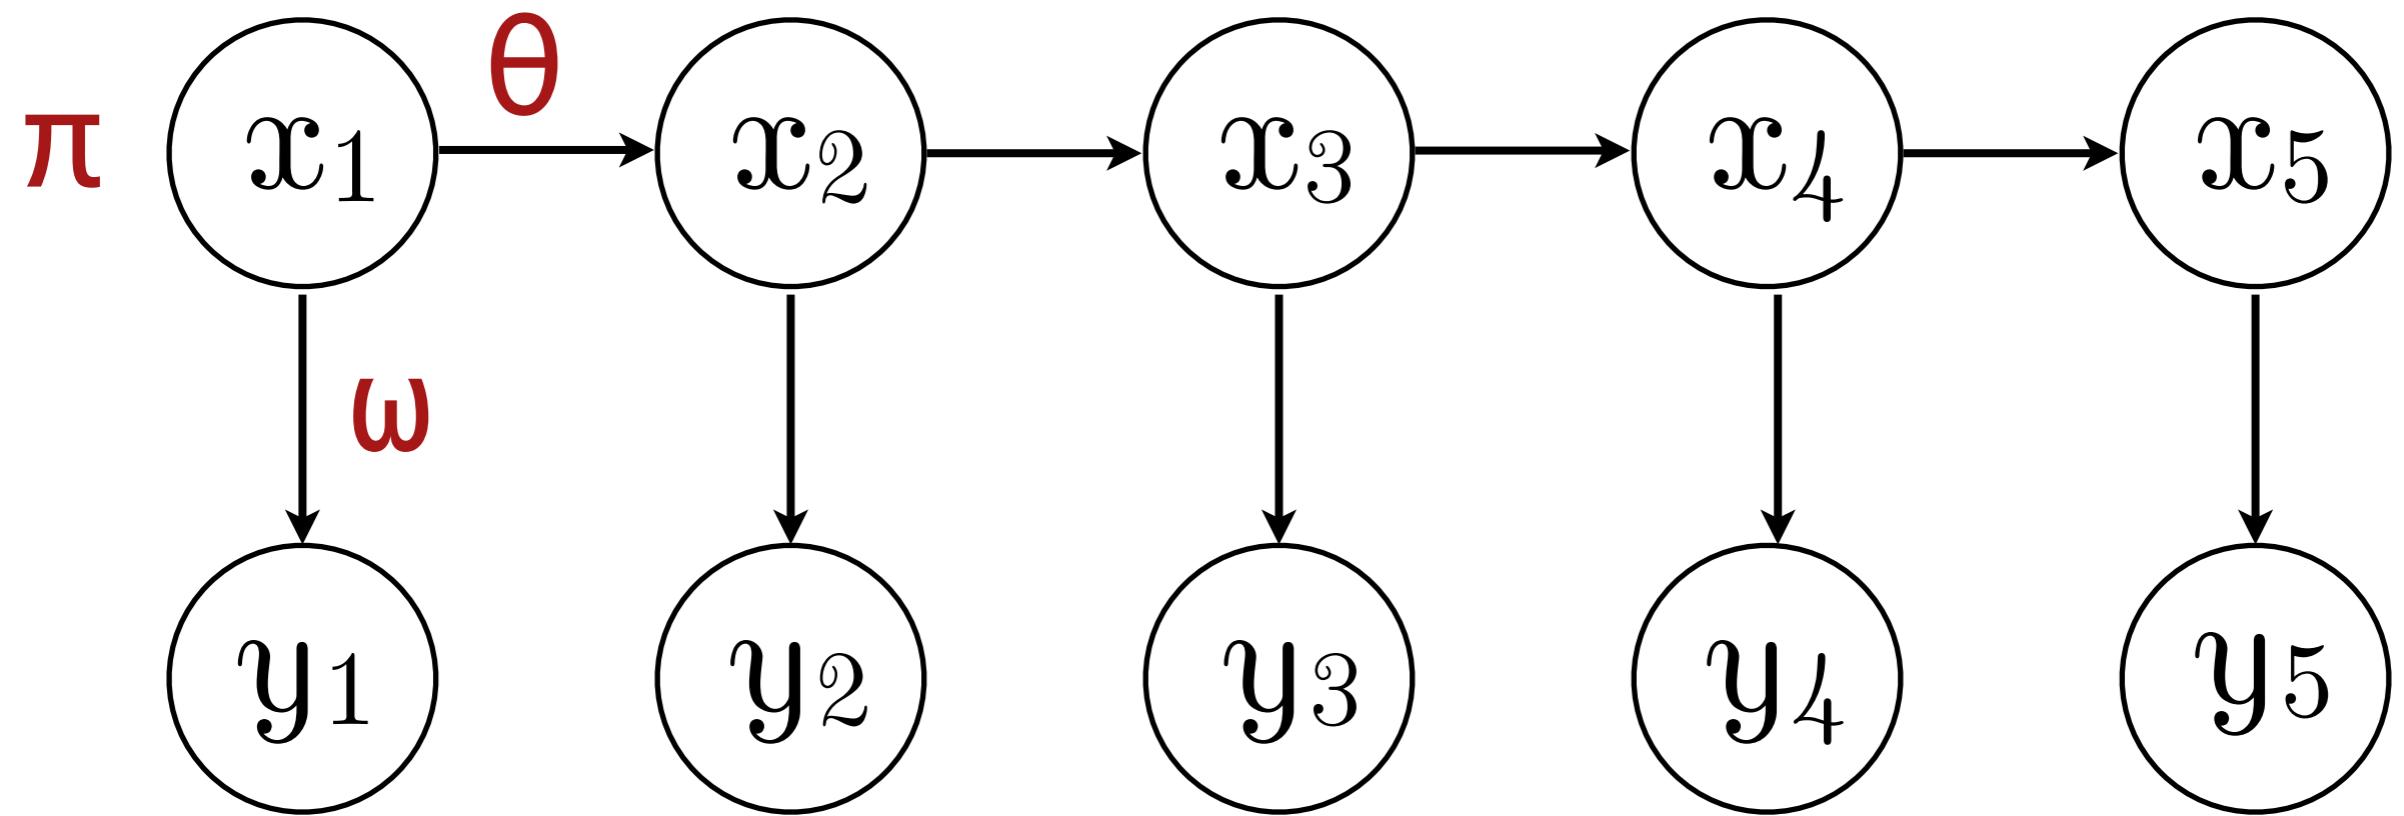
... that aren't
optimized

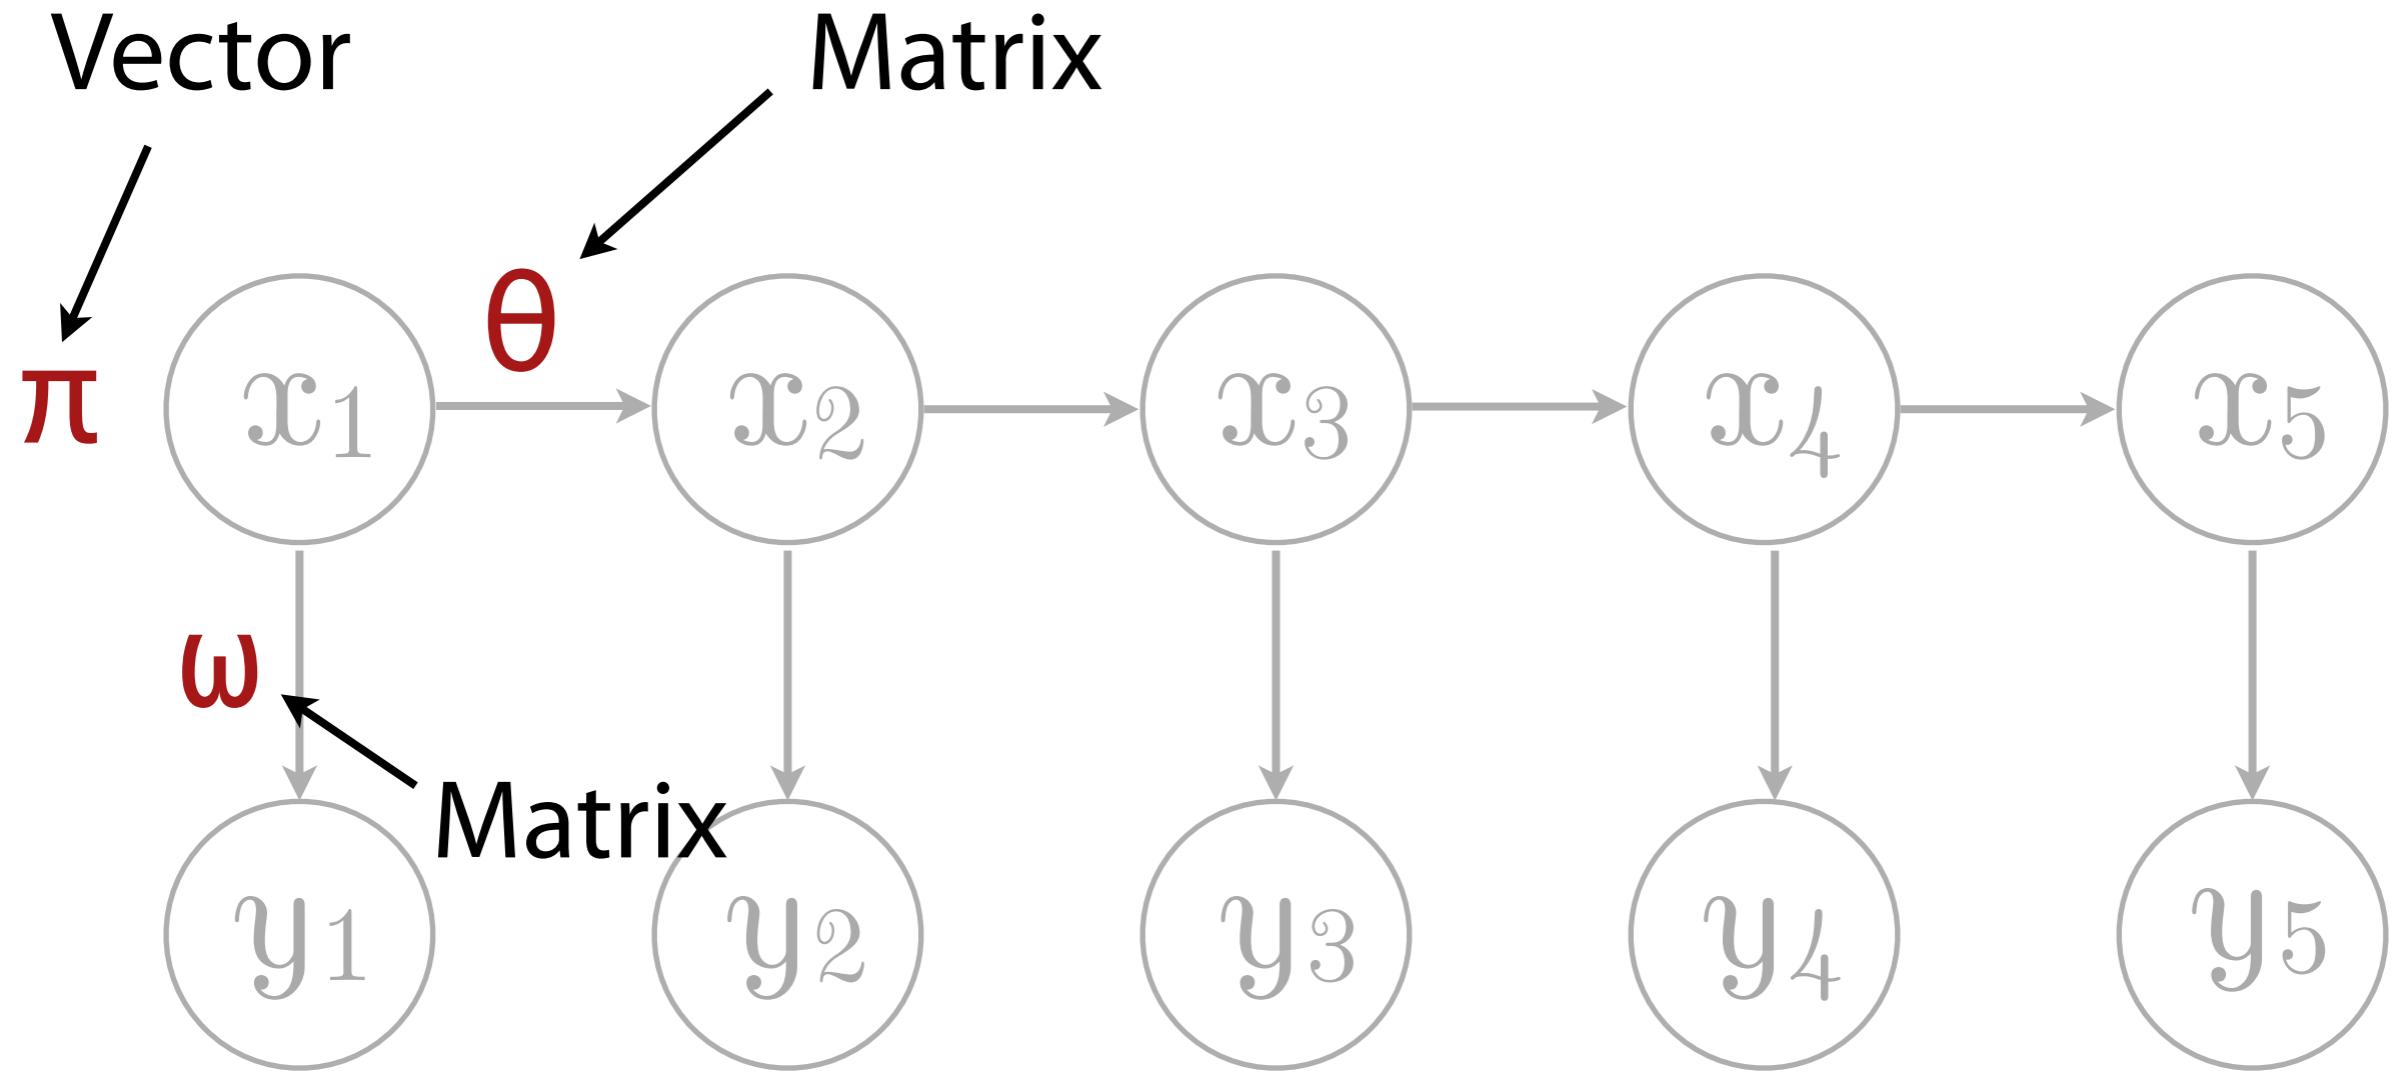
unicode identifiers

array copying gotchas

simple and easy
parallel programming







How do we learn these?

How do we learn these?

$$\alpha_t(x_t) =: p(x_t, y_0, \dots, y_t)$$

Seed: $\alpha_0(x_0) = p(y_0|x_0) \pi_{x_0}$

Iterate: $\alpha_t(x_t) = p(y_t|x_t) \sum_{x_{t-1}} p(x_t|x_{t-1}) \alpha_{t-1}(x_{t-1})$

$$p(x_t|y_0, \dots, y_t) = \frac{\alpha_t(x_t)}{\sum_{x'_t} \alpha_t(x'_t)}$$

How do we learn these?

$$\beta_t(x_t) = p(y_{t+1}, \dots, y_T | x_t)$$

Seed: $\beta_T(x_T) = 1$

Iterate: $\beta_t(x_t) = \sum_{x_{t+1}} p(x_{t+1}|x_t) p(y_{t+1}|x_{t+1}) \beta_{t+1}(x_{t+1})$

$$p(x_t|y_0, \dots, y_T) = \frac{\alpha_t(x_t) \beta_t(x_t)}{\sum_{x'_t} \alpha_t(x'_t) \beta_t(x'_t)} =: \gamma_t(x_t)$$

How do we learn these?

$$\begin{aligned} p(x_t, x_{t+1} | y_0, \dots, y_T) &= \frac{1}{Z} \alpha_t(x_t) p(x_{t+1} | x_t) p(y_{t+1} | x_{t+1}) \beta_{t+1}(x_{t+1}) = \\ &\frac{\gamma_t(x_t)}{\beta_t(x_t)} p(x_{t+1} | x_t) p(y_{t+1} | x_{t+1}) \beta_{t+1}(x_{t+1}) =: \xi_t(x_t, x_{t+1}) \end{aligned}$$

How do we learn these?

$$\pi_i^{(i+1)} = \gamma_0(i)$$

$$\omega_{i,j}^{(i+1)} = \frac{\sum_{t: y_t=j} \gamma_t(i)}{\sum_{t=0}^T \gamma_t(i)}$$

$$\theta_{i,j}^{(i+1)} = \frac{\sum_{t=0}^{T-1} \xi_t(i, j)}{\sum_{t=0}^{T-1} \gamma_t(i)},$$

it's hard

it's hard

(no matter what tools you use)

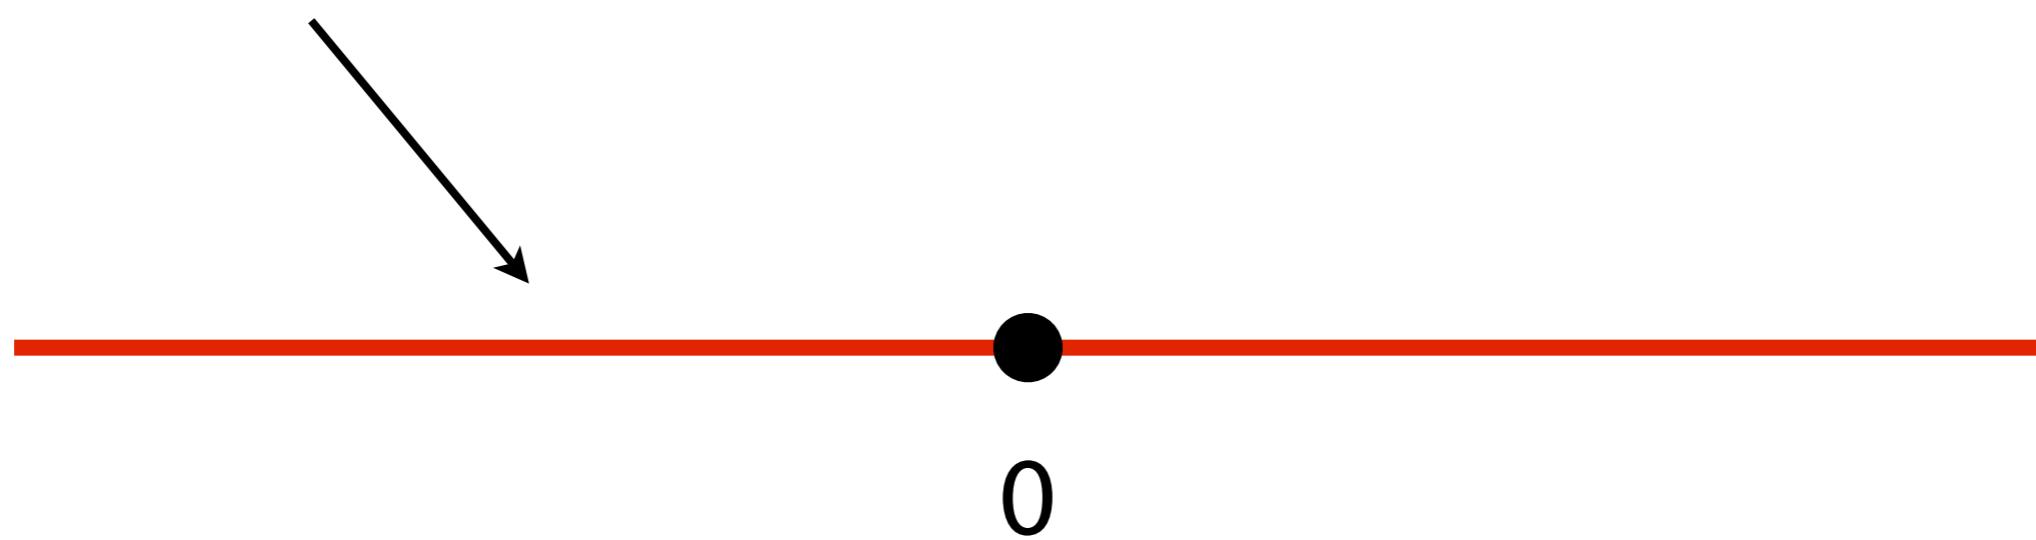
```
type HMM{T <: FloatingPoint}
    nh::Int
    no::Int
    t::Int
    π::Vector{T}
    ω::Matrix{T}
    θ::Matrix{T}
end
```

```
type HMM[T <: FloatingPoint]
    nh::Int
    no::Int
    t::Int
    π::Vector{T}
    ω::Matrix{T}
    θ::Matrix{T}
end
```

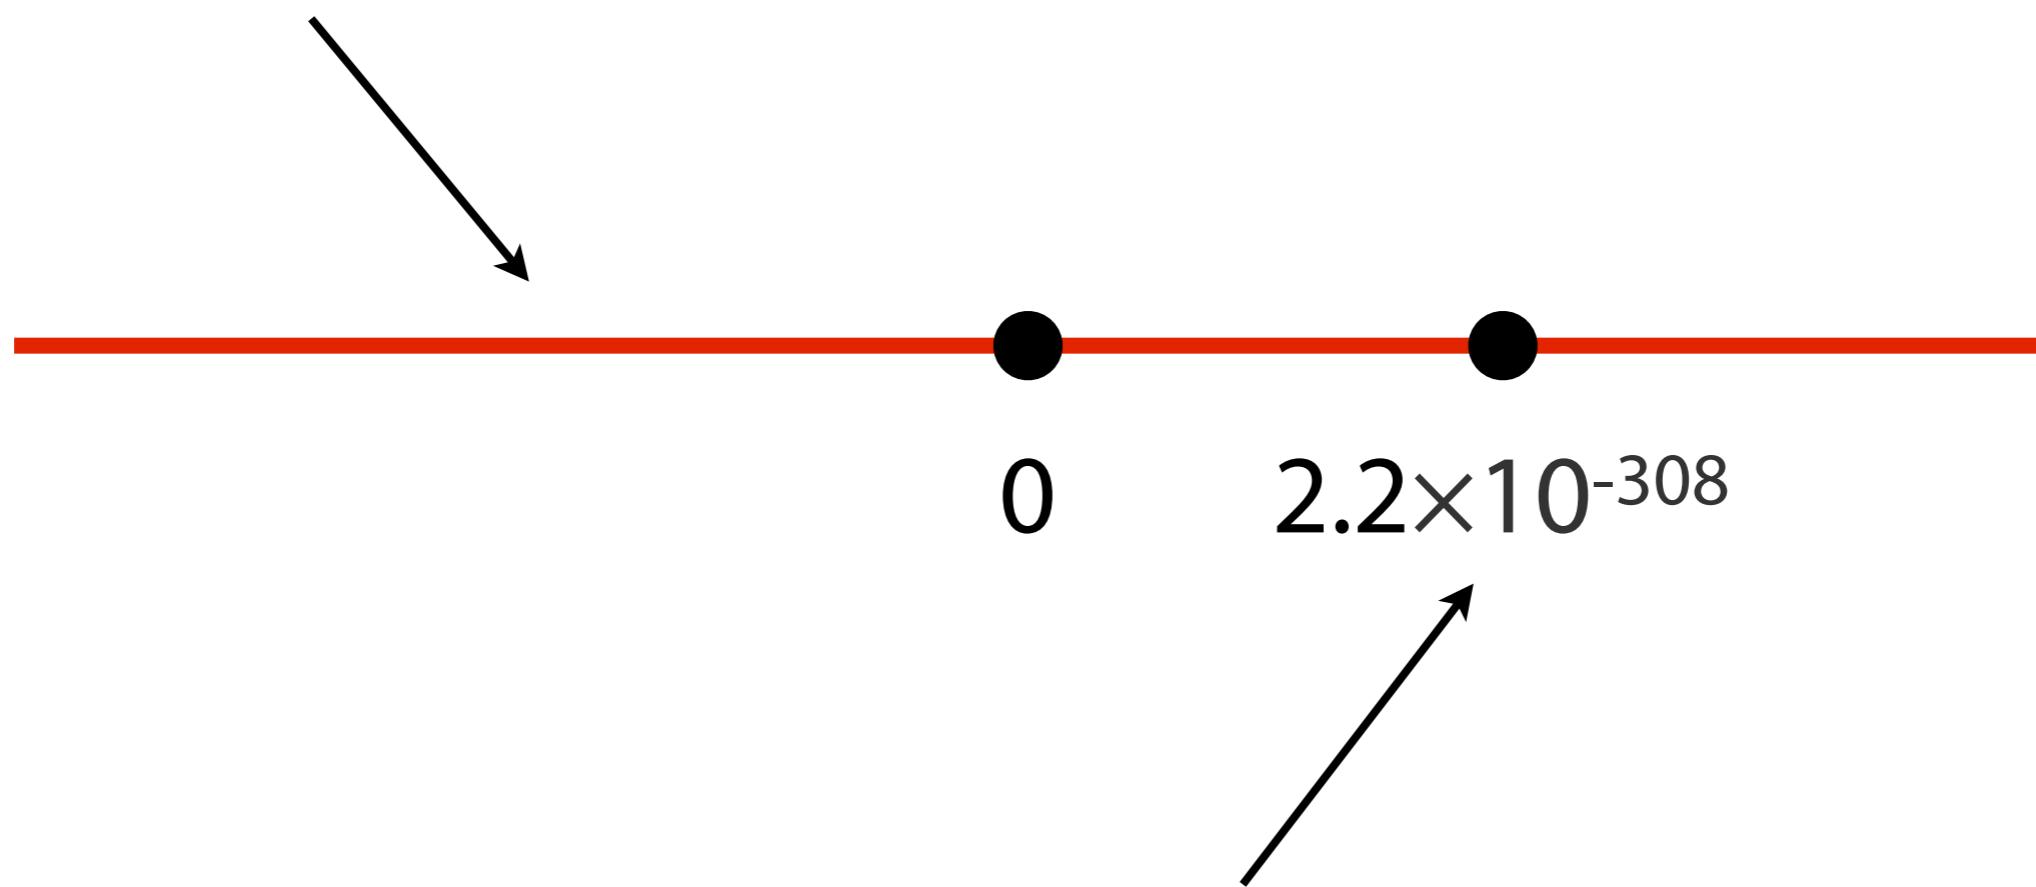
underflow

IEEE 754

the real line

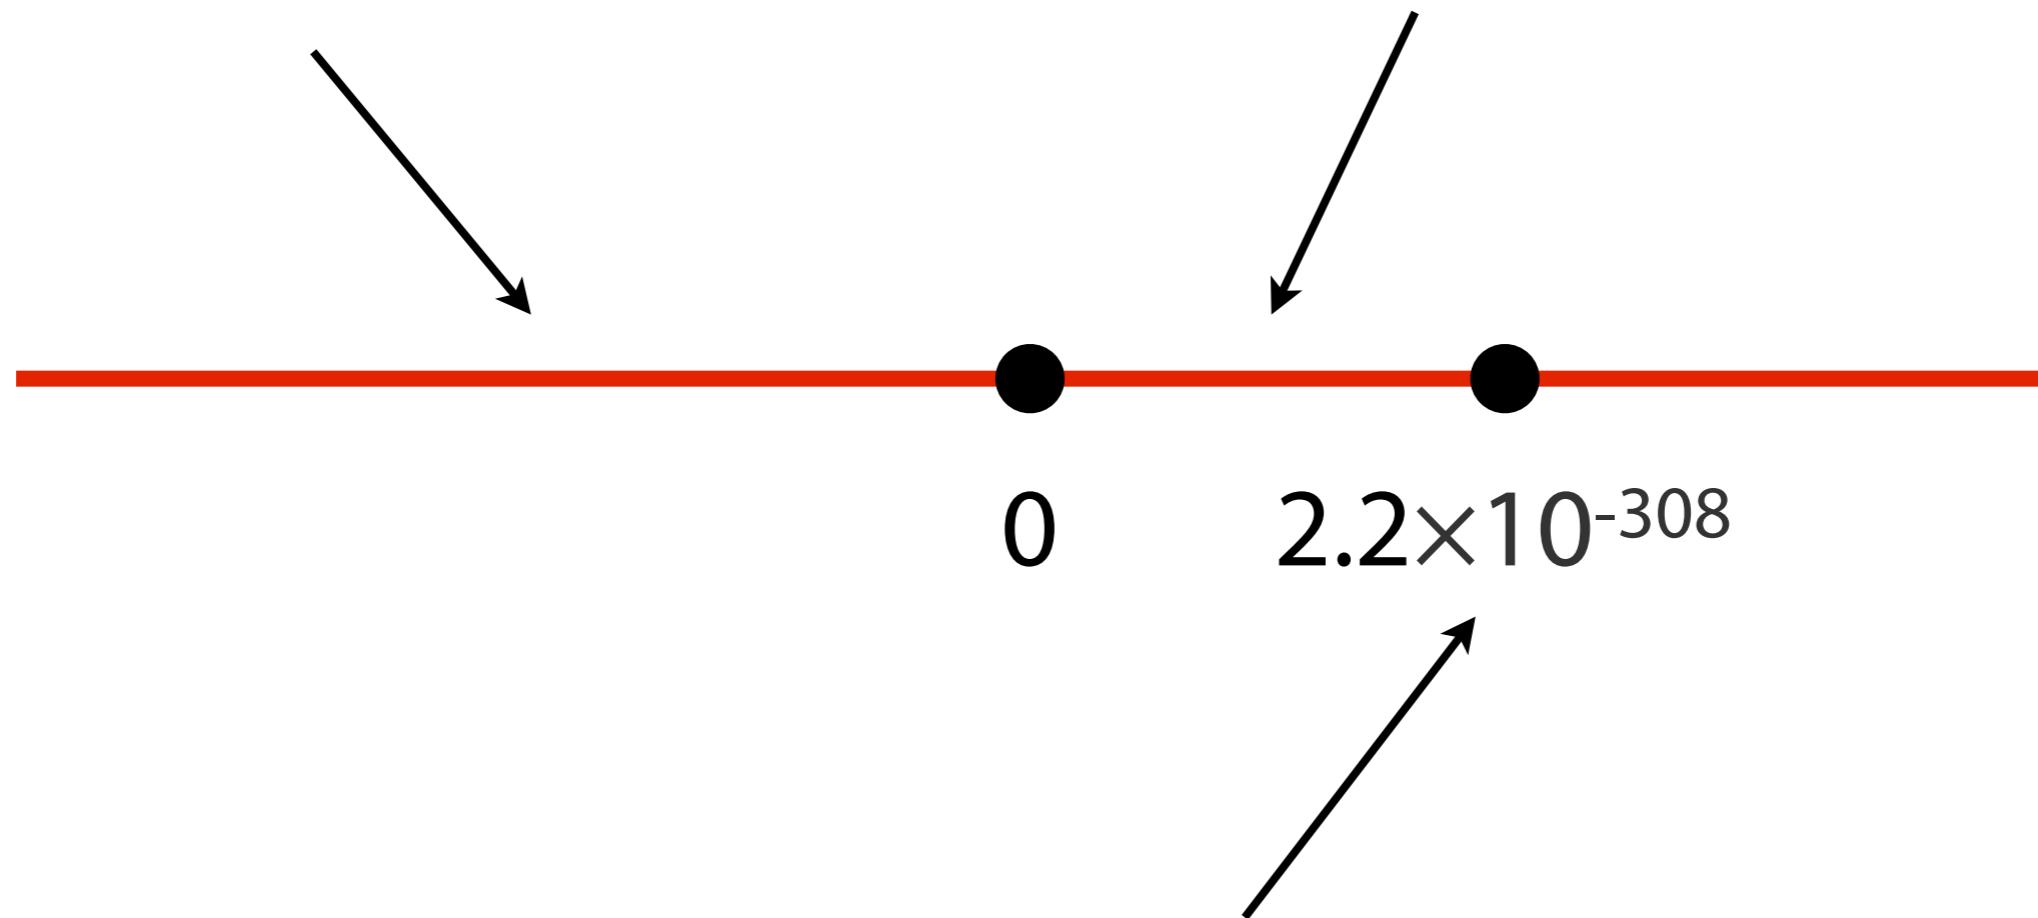


the real line



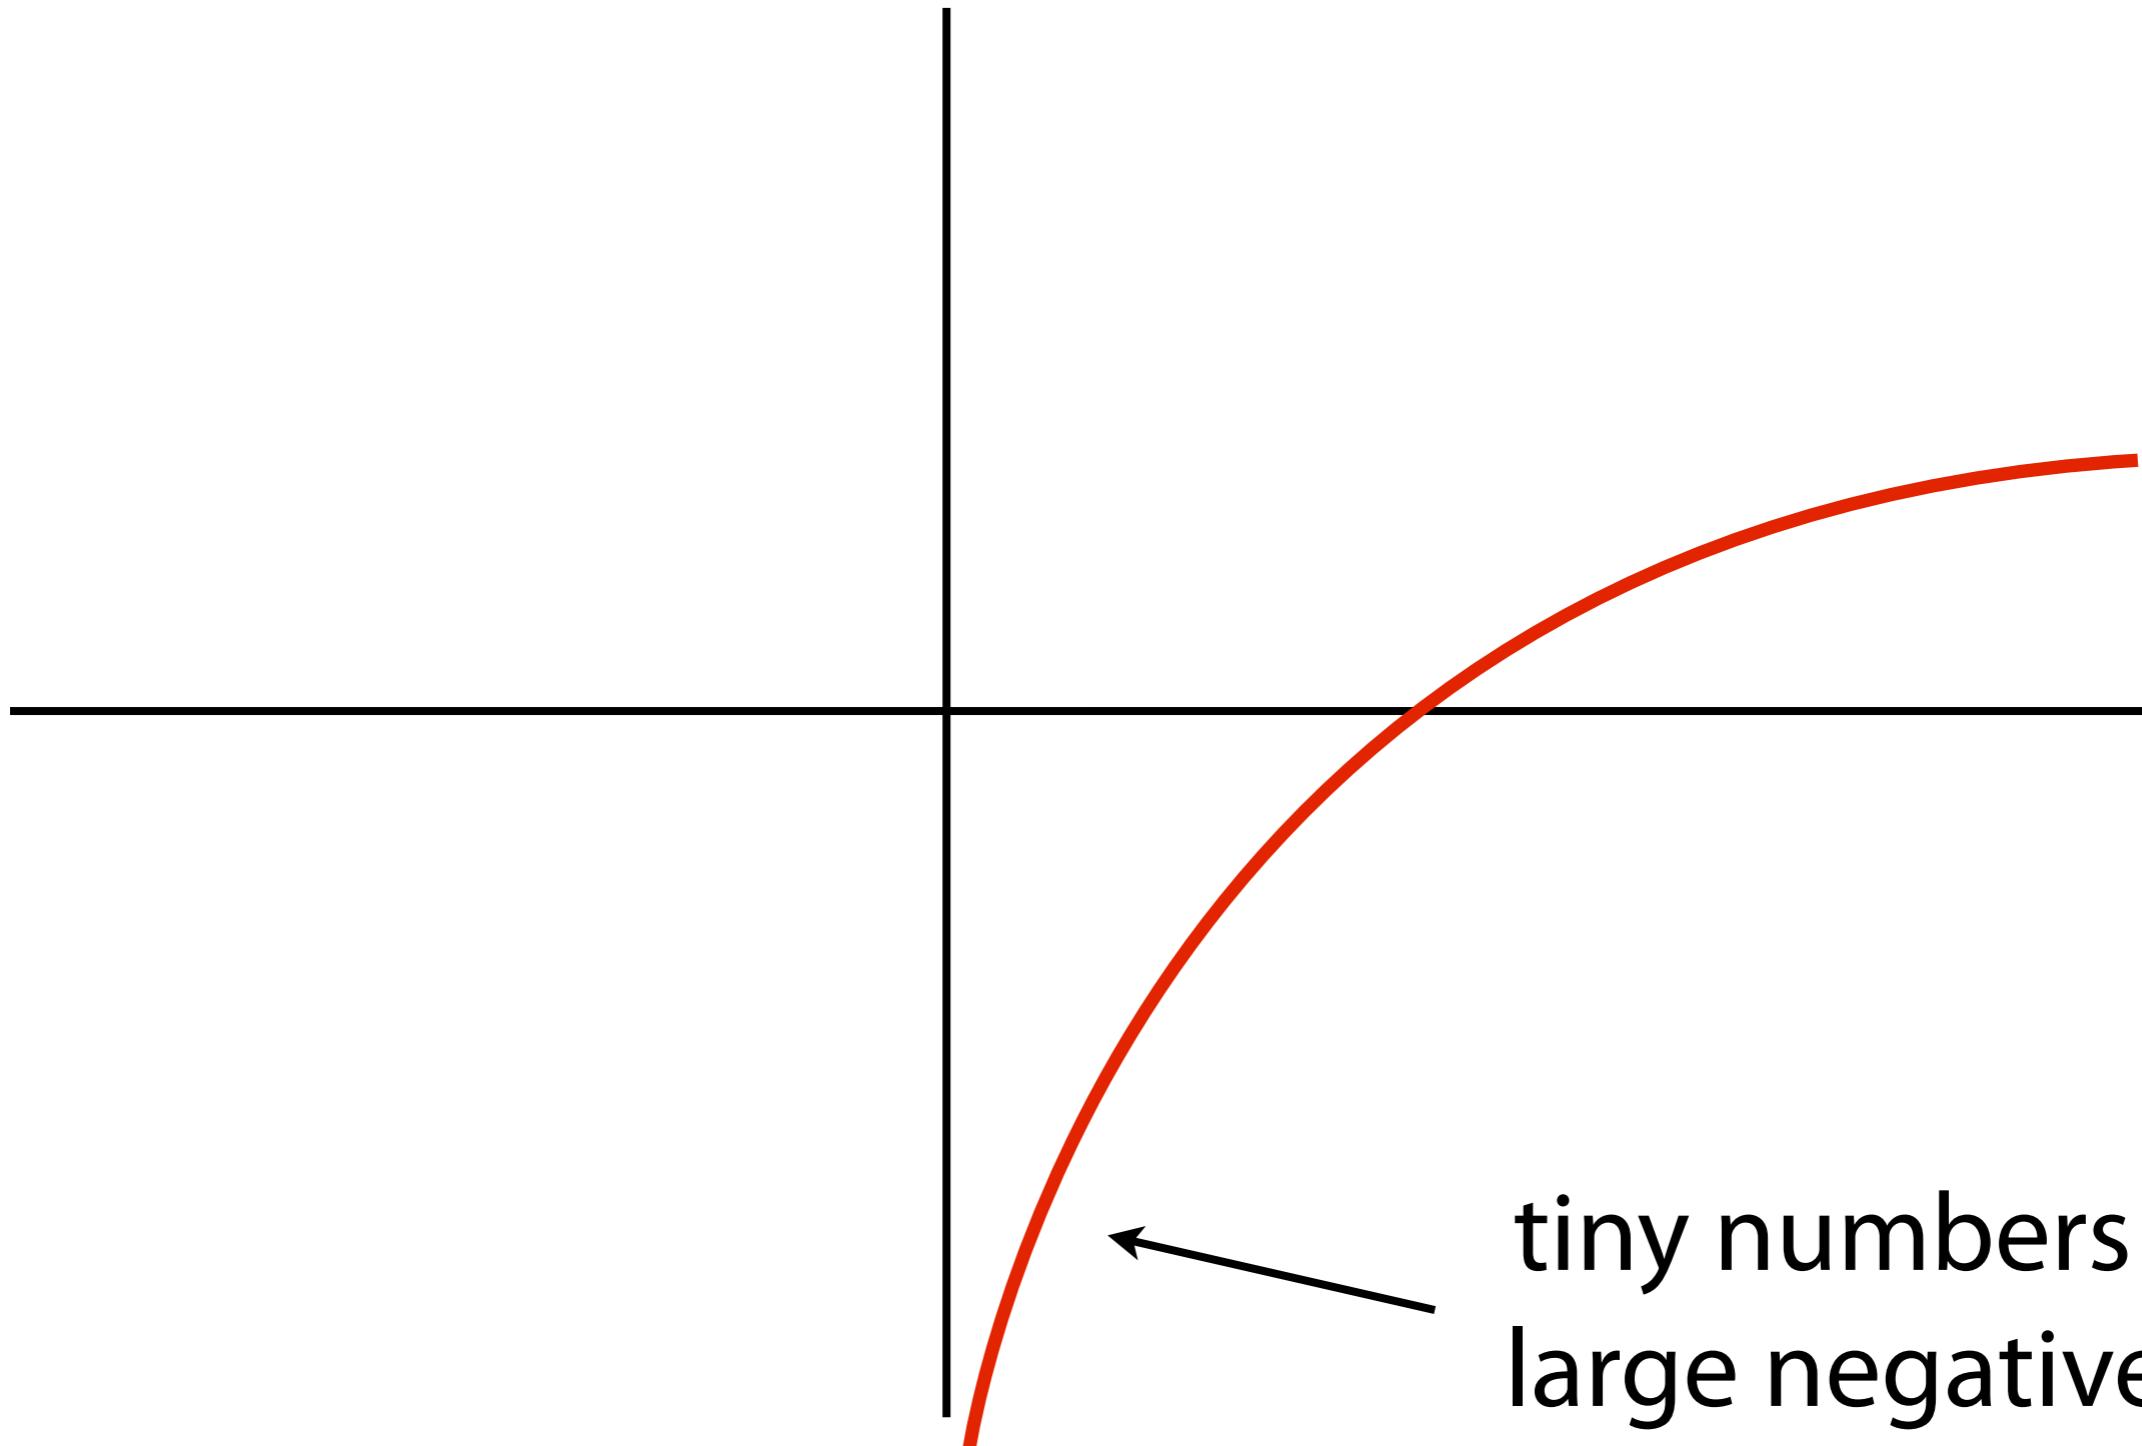
smallest number representable
by 64-bit floats

the real line



smallest number representable
by 64-bit floats

use logs



use logs

$$\log(xy) = \log x + \log y$$

use logs

$$\log(xy) = \log x + \log y$$

$$\log(x + y) =$$

$$\log x + \log(1 + e^{\log y - \log x})$$

```
immutable LogFloat <: FloatingPoint
    n::Float64
end
```

```
Base.convert(::Type{LogFloat},  
           f::Float64) = LogFloat(log(f))
```

```
Base.convert(::Type{LogFloat},  
           f::Float64) = LogFloat(log(f))
```

```
Base.convert(::Type{Float64},  
           lf::LogFloat) = exp(lf.n)
```

```
Base.convert(::Type{LogFloat},  
           f::Float64) = LogFloat(log(f))
```

```
Base.convert(::Type{Float64},  
           lf::LogFloat) = exp(lf.n)
```

```
Base.promote_rule(::Type{LogFloat},  
                  ::Type{Float64}) = LogFloat
```

```
# a LogFloat(n) represents e^n
function Base.show(io::IO, lf::LogFloat)
    n = lf.n
    print(io, "e^$n")
end
```

```
function (*) (lf1::LogFloat, lf2::LogFloat)
    return LogFloat(lf1.n + lf2.n)
end
```

```
function (*) (lf1::LogFloat, lf2::LogFloat)
    return LogFloat(lf1.n + lf2.n)
end
```

```
function (/) (lf1::LogFloat, lf2::LogFloat)
    return LogFloat(lf1.n - lf2.n)
end
```

```
function (*) (lf1::LogFloat, lf2::LogFloat)
    return LogFloat(lf1.n + lf2.n)
end
```

```
function (/) (lf1::LogFloat, lf2::LogFloat)
    return LogFloat(lf1.n - lf2.n)
end
```

```
function (+) (lf1::LogFloat, lf2::LogFloat)
    if lf1.n > lf2.n
        x = lf1.n
        y = lf2.n
    else
        x = lf2.n
        y = lf1.n
    end
    LogFloat(x + log1p(exp(y-x)))
end
```

```
function (<)(lf1::LogFloat,lf2::LogFloat)
    return lf1.n < lf2.n
end
```

$$\frac{\alpha_t(x_t) \beta_t(x_t)}{\sum_{x'_t} \alpha_t(x'_t) \beta_t(x'_t)} =: \gamma_t(x_t)$$

```
# gamma[i,t] = state i ,time t
gamma(α,β) = broadcast(/, α .* β, sum(α .* β,1))
```

$$\frac{\alpha_t(x_t) \beta_t(x_t)}{\sum_{x'_t} \alpha_t(x'_t) \beta_t(x'_t)} =: \gamma_t(x_t)$$

```
# gamma[i,t] = state i ,time t  
gamma(α,β) = broadcast(/, α .* β, sum(α .* β,1))
```

$$\frac{\alpha_t(x_t) \beta_t(x_t)}{\sum_{x'_t} \alpha_t(x'_t) \beta_t(x'_t)} =: \gamma_t(x_t)$$

```
# gamma[i,t] = state i ,time t  
gamma(α,β) = broadcast(/, α .* β, sum(α .* β,1))
```

$$\frac{\alpha_t(x_t) \beta_t(x_t)}{\sum_{x'_t} \alpha_t(x'_t) \beta_t(x'_t)} =: \gamma_t(x_t)$$

```
# gamma[i,t] = state i ,time t  
gamma(α,β) = broadcast(/, α .* β, sum(α .* β,1))
```

$$\frac{\alpha_t(x_t) \beta_t(x_t)}{\sum_{x'_t} \alpha_t(x'_t) \beta_t(x'_t)} =: \gamma_t(x_t)$$

```
# gamma[i,t] = state i ,time t  
gamma(α,β) = broadcast(/, α .* β, sum(α .* β,1))
```

$$\frac{\alpha_t(x_t) \beta_t(x_t)}{\sum_{x'_t} \alpha_t(x'_t) \beta_t(x'_t)} =: \gamma_t(x_t)$$

```
# gamma[i,t] = state i ,time t
gamma(α,β) = broadcast(/, α .* β, sum(α .* β,1))
```

```

# gamma[i,t] = state i ,time t
function gamma( $\alpha$ , $\beta$ )
    ret =  $\alpha$  .*  $\beta$ 
    # sums[t] is sum at time t over all states
    sums = sum(ret,1)
    for i in 1:size(ret,1)
        for j in 1:size(ret,2)
            ret[i,j] = ret[i,j] / sums[j]
        end
    end
    return ret
end

```

```
hmm = HMM{LogFloat}(20, 27, 1000)
```

```
train!(text, hmm)
```

```
correct(hmm, corrupted_text)
```

```
sample(hmm, n)
```

THE
GOOD

THE
BAD

type system /
numeric promotions

THE GOOD

THE BAD

type system /
numeric promotions

easy to map from
math onto code
(readability)

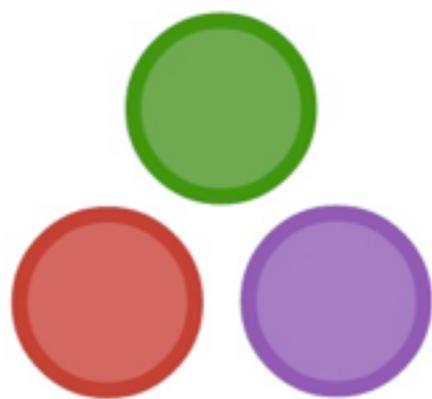
THE GOOD

THE BAD

type system /
numeric promotions

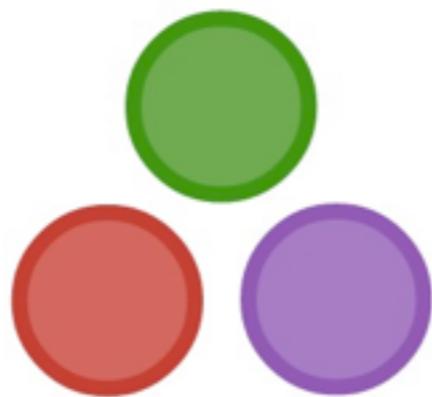
no debugger

easy to map from
math onto code
(readability)



THE
GOOD

THE
BAD



THE
GOOD > THE
BAD

questions?



@porterjamesj

github.com/porterjamesj/254

porterjamesj@gmail.com